Week 2 – Python Bootcamp

02613 Python and High-Performance Computing

Last week in 02613…

Corrected/updated exam date

- Exam on 1 June 2026
  - DTU's websites (student.dtu.dk) authoritative
  - DTU's calendar for Spring semester 2026 updated last week

Week 1 lecture recording not available

VPN issues

- Upgrade to the latest client version https://net.ait.dtu.dk/vpn
- If that does not help, contact IT itservice@dtu.dk.
- Alternative: SSH keys https://www.hpc.dtu.dk/?page_id=4317
  - Setup once on a DTU network or on VPN
  - From there, no VPN required
- “Emergency” alternative: https://remote.dtu.dk
  - Windows desktop ”within DTUs walls”
  - No VPN required (but maybe need 2FA)

Login nodes

Four login nodes available

ssh <username>@login.hpc.dtu.dk
ssh <username>@login2.hpc.dtu.dk
ssh <username>@login.gbar.dtu.dk
ssh <username>@login2.gbar.dtu.dk

Login nodes

Remember:

Login nodes are *NOT* for computations

Please keep them free to handle logins

Use **linuxsh** for an interactive session (test, debug)

Use **bsub** for “real” computations

Login nodes and VS Code

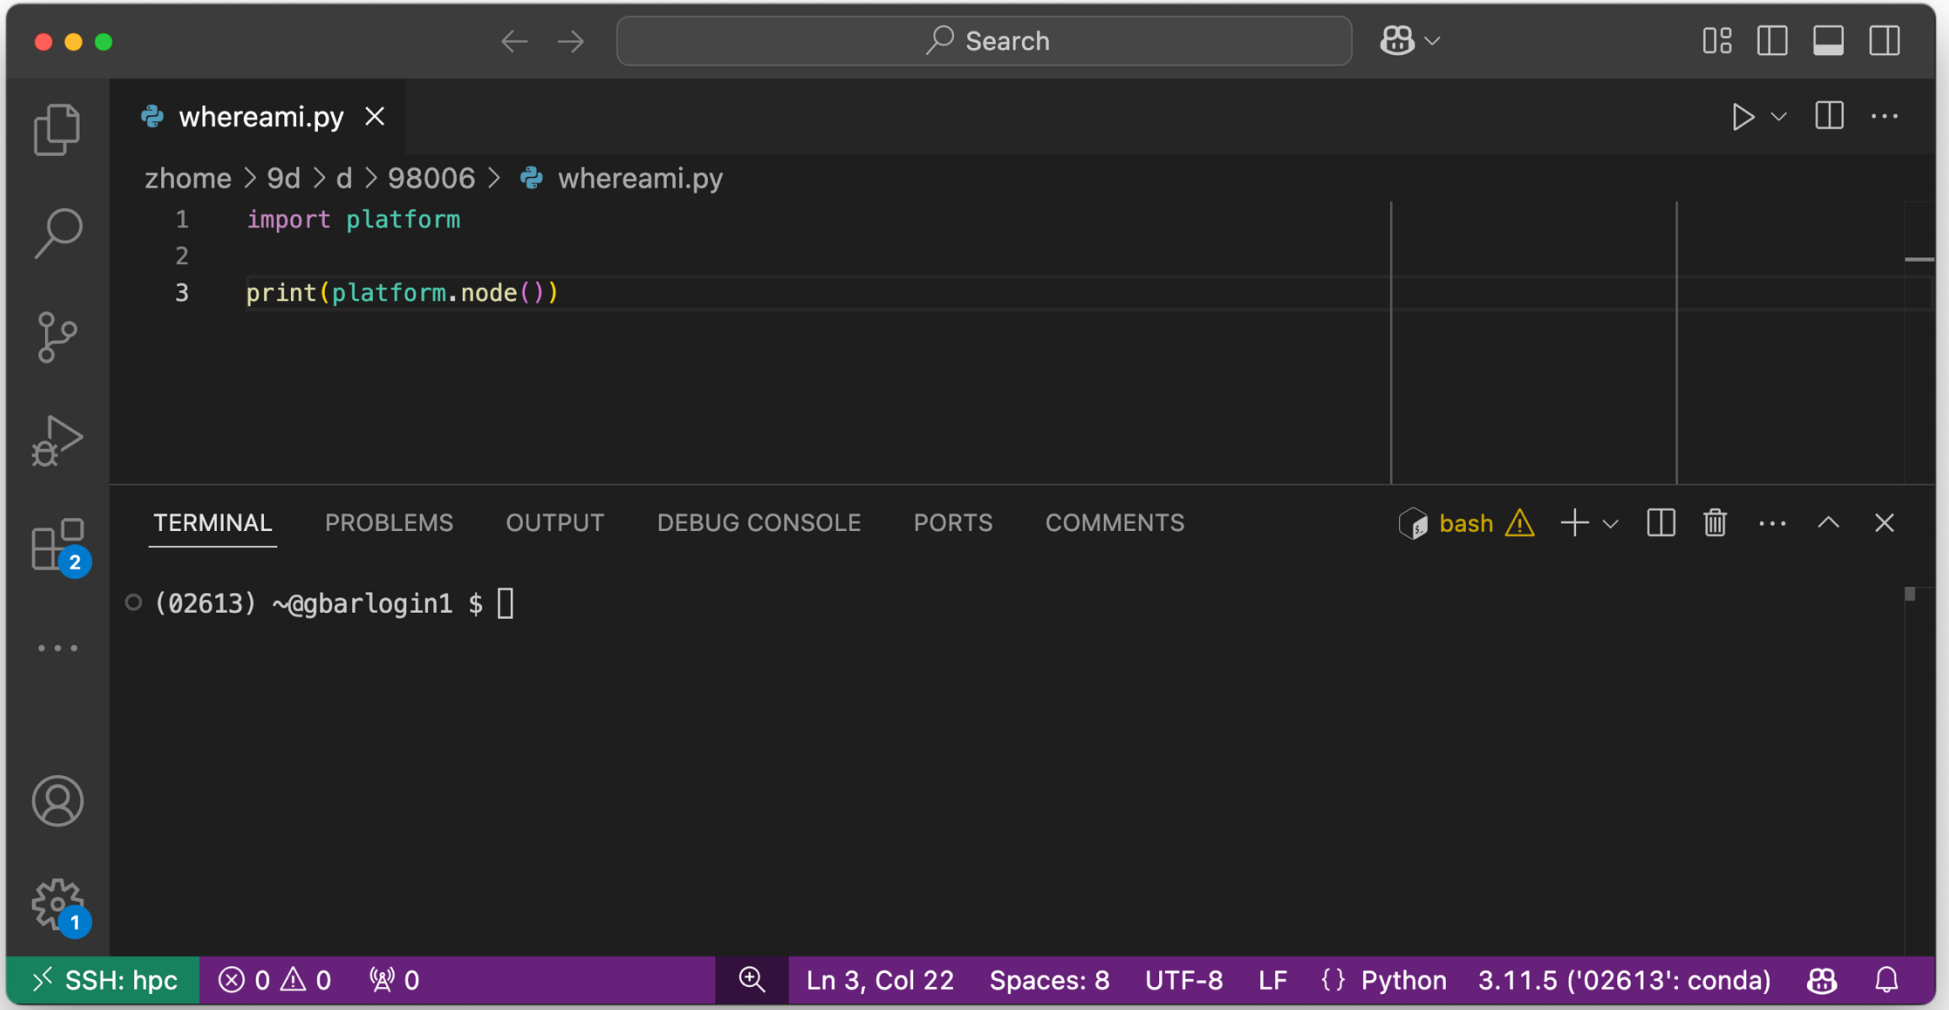

Login nodes and VS Code

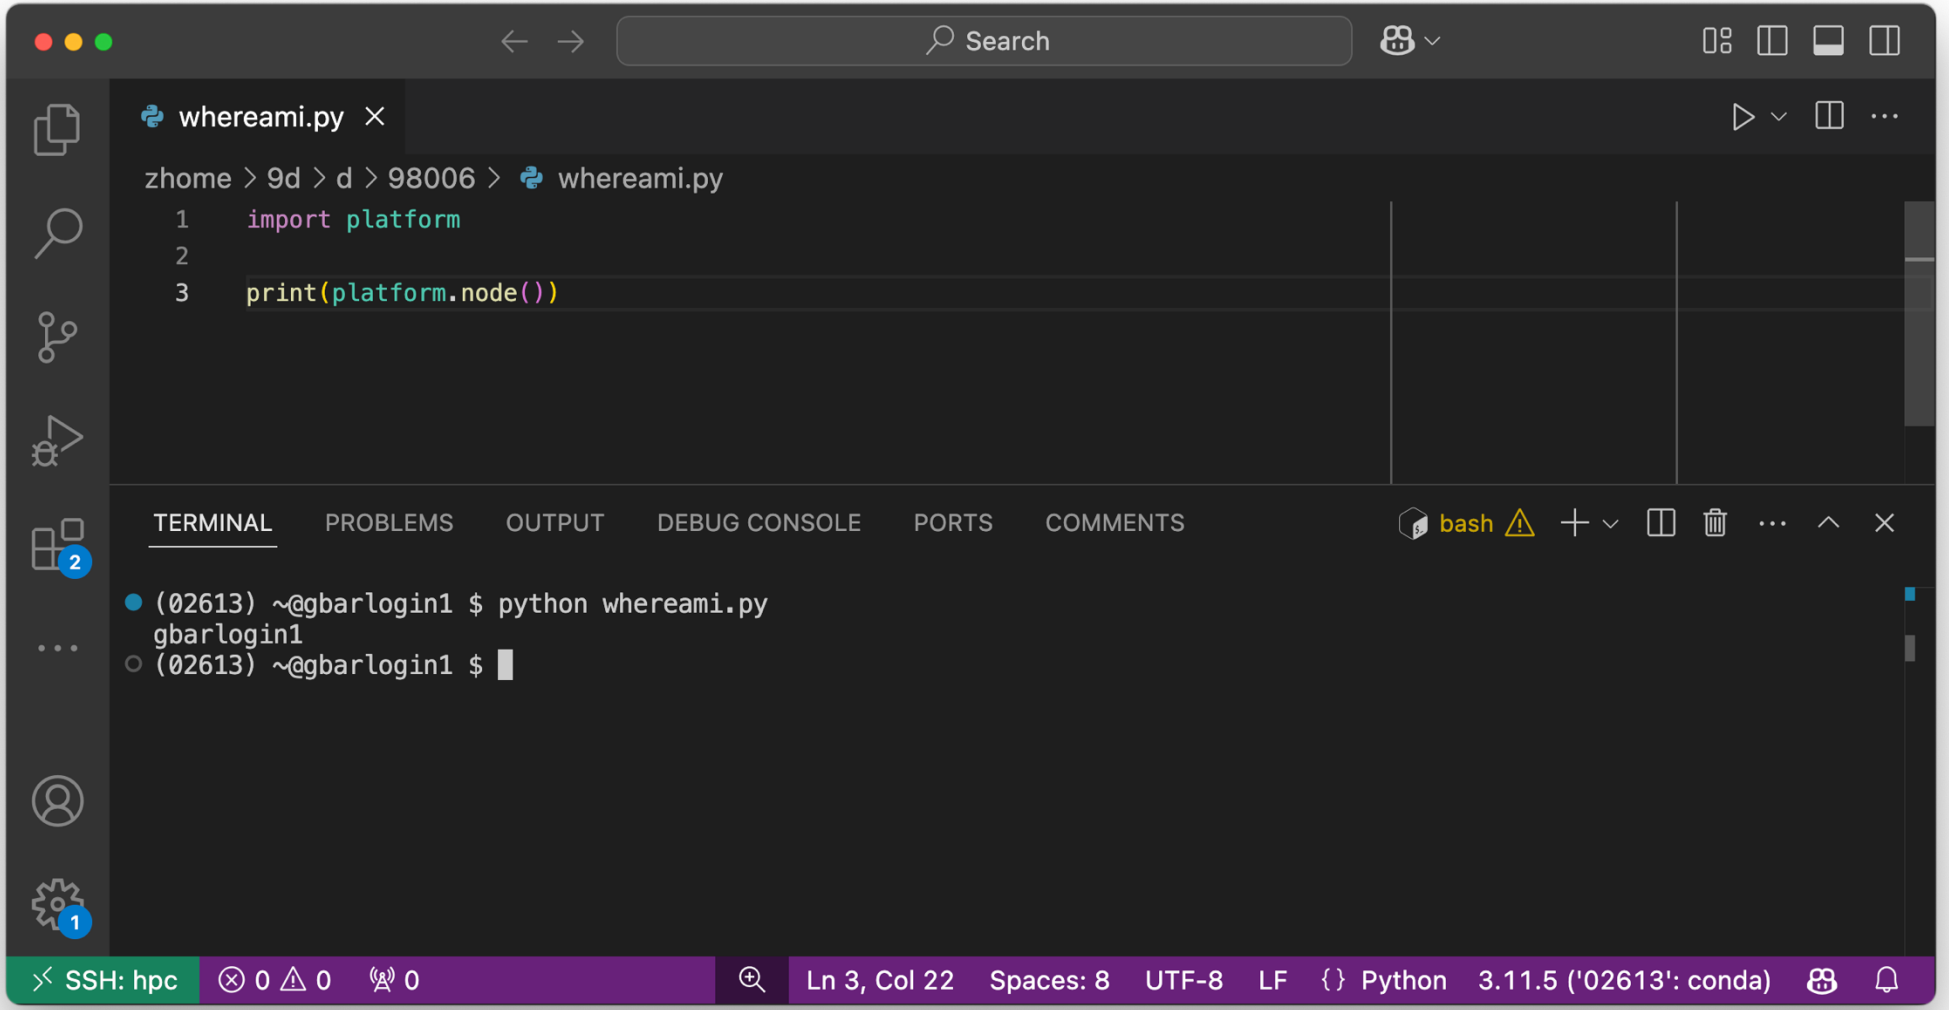

Login nodes and VS Code

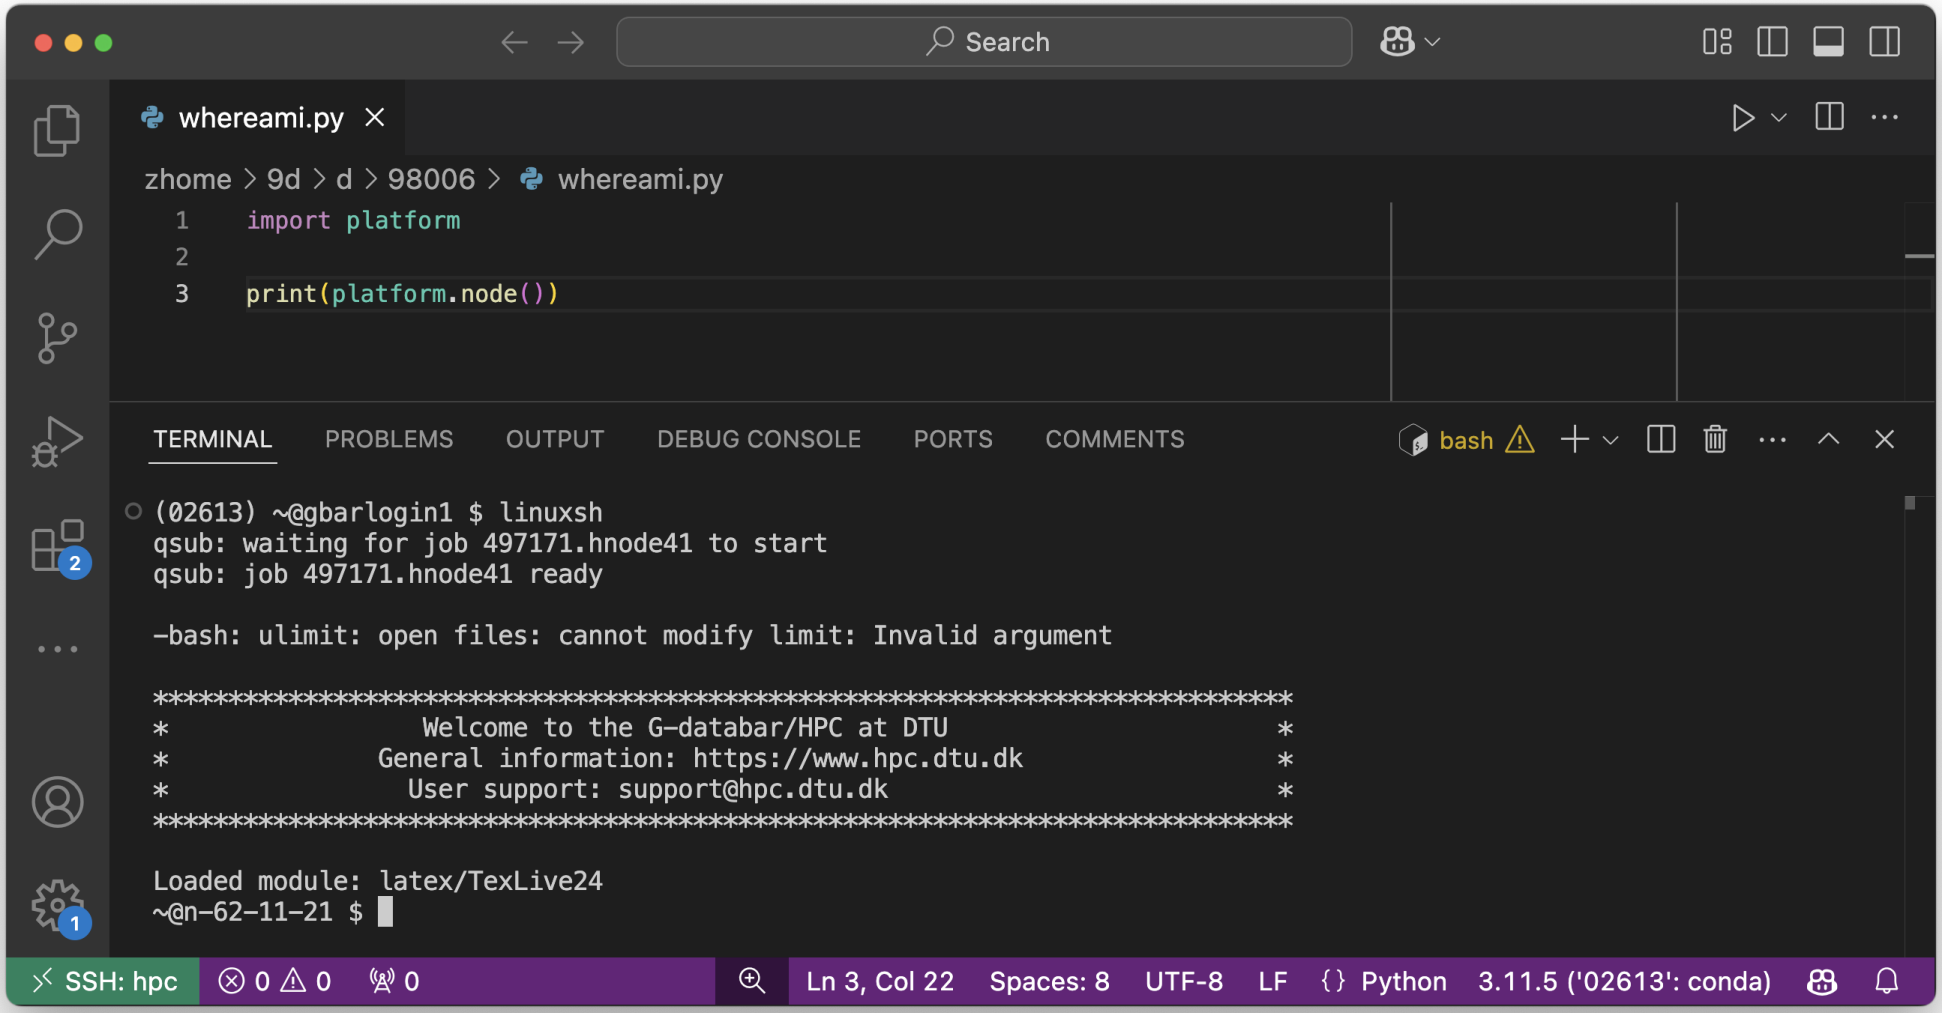

Login nodes and VS Code

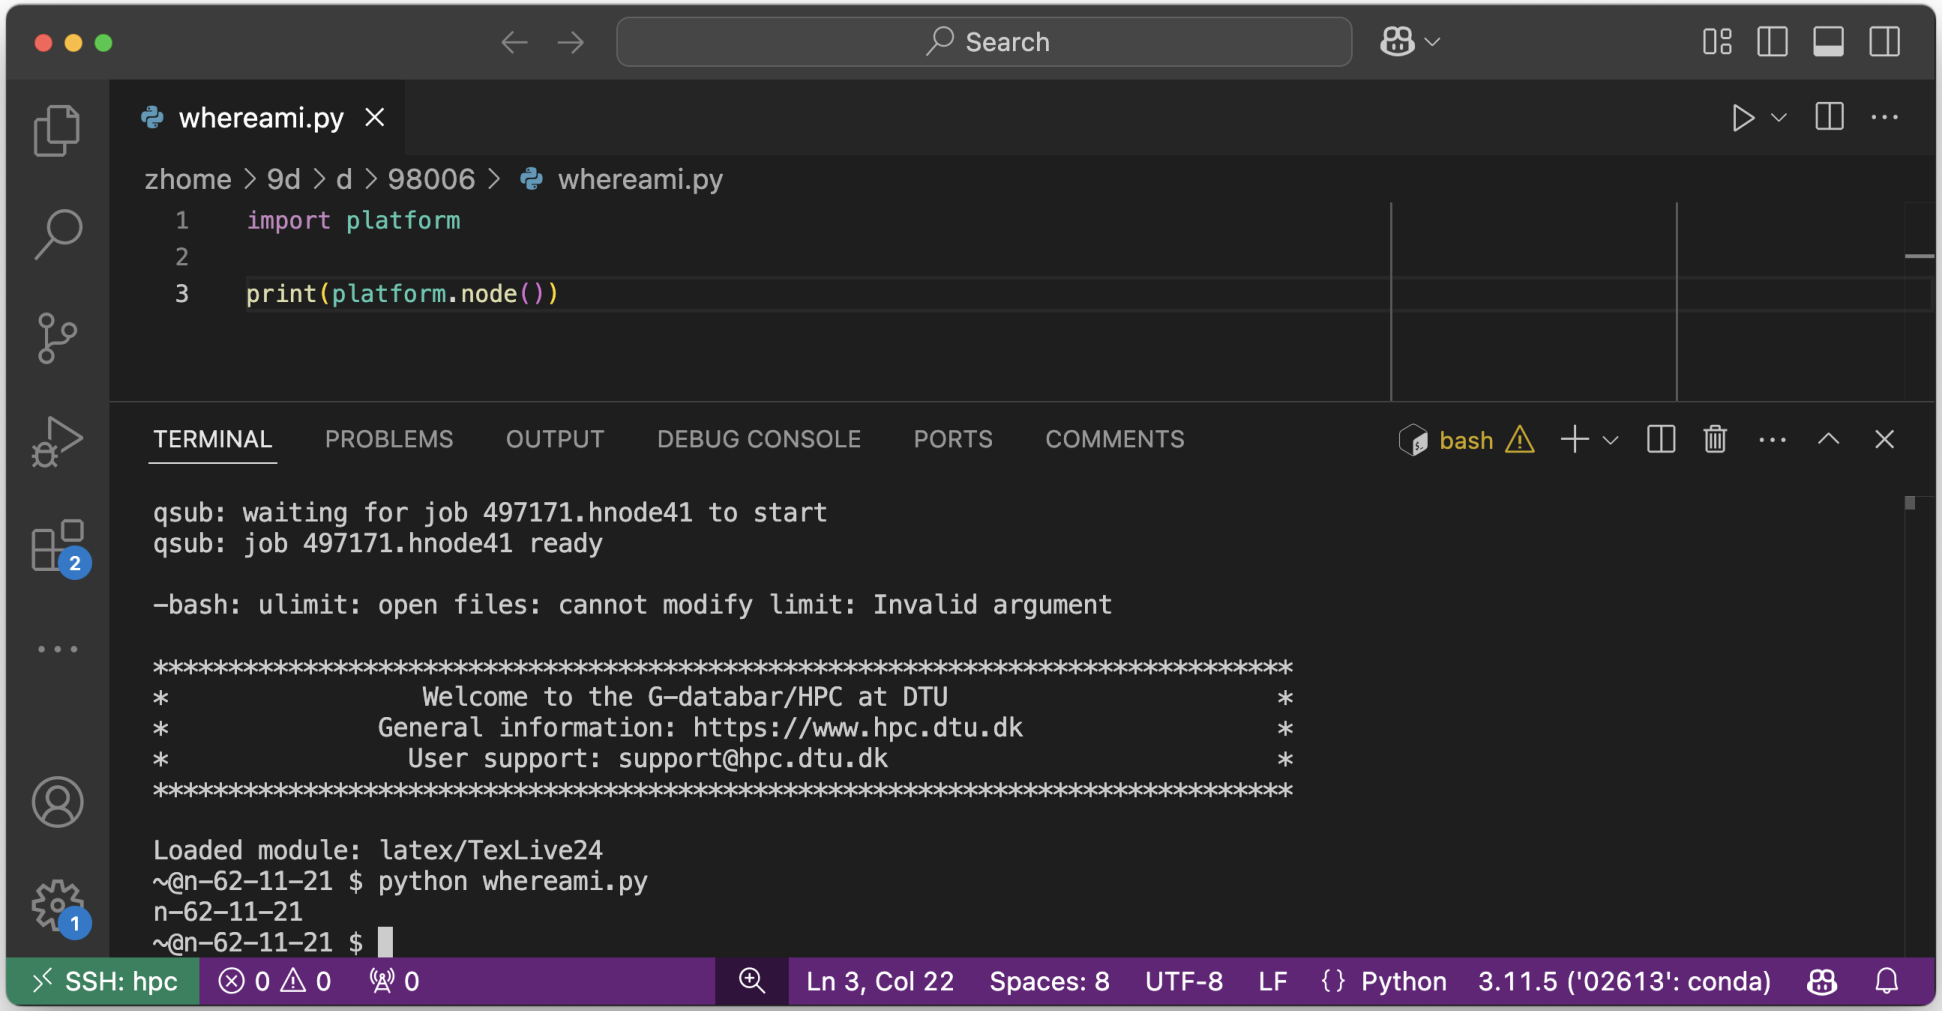

Login nodes and VS Code

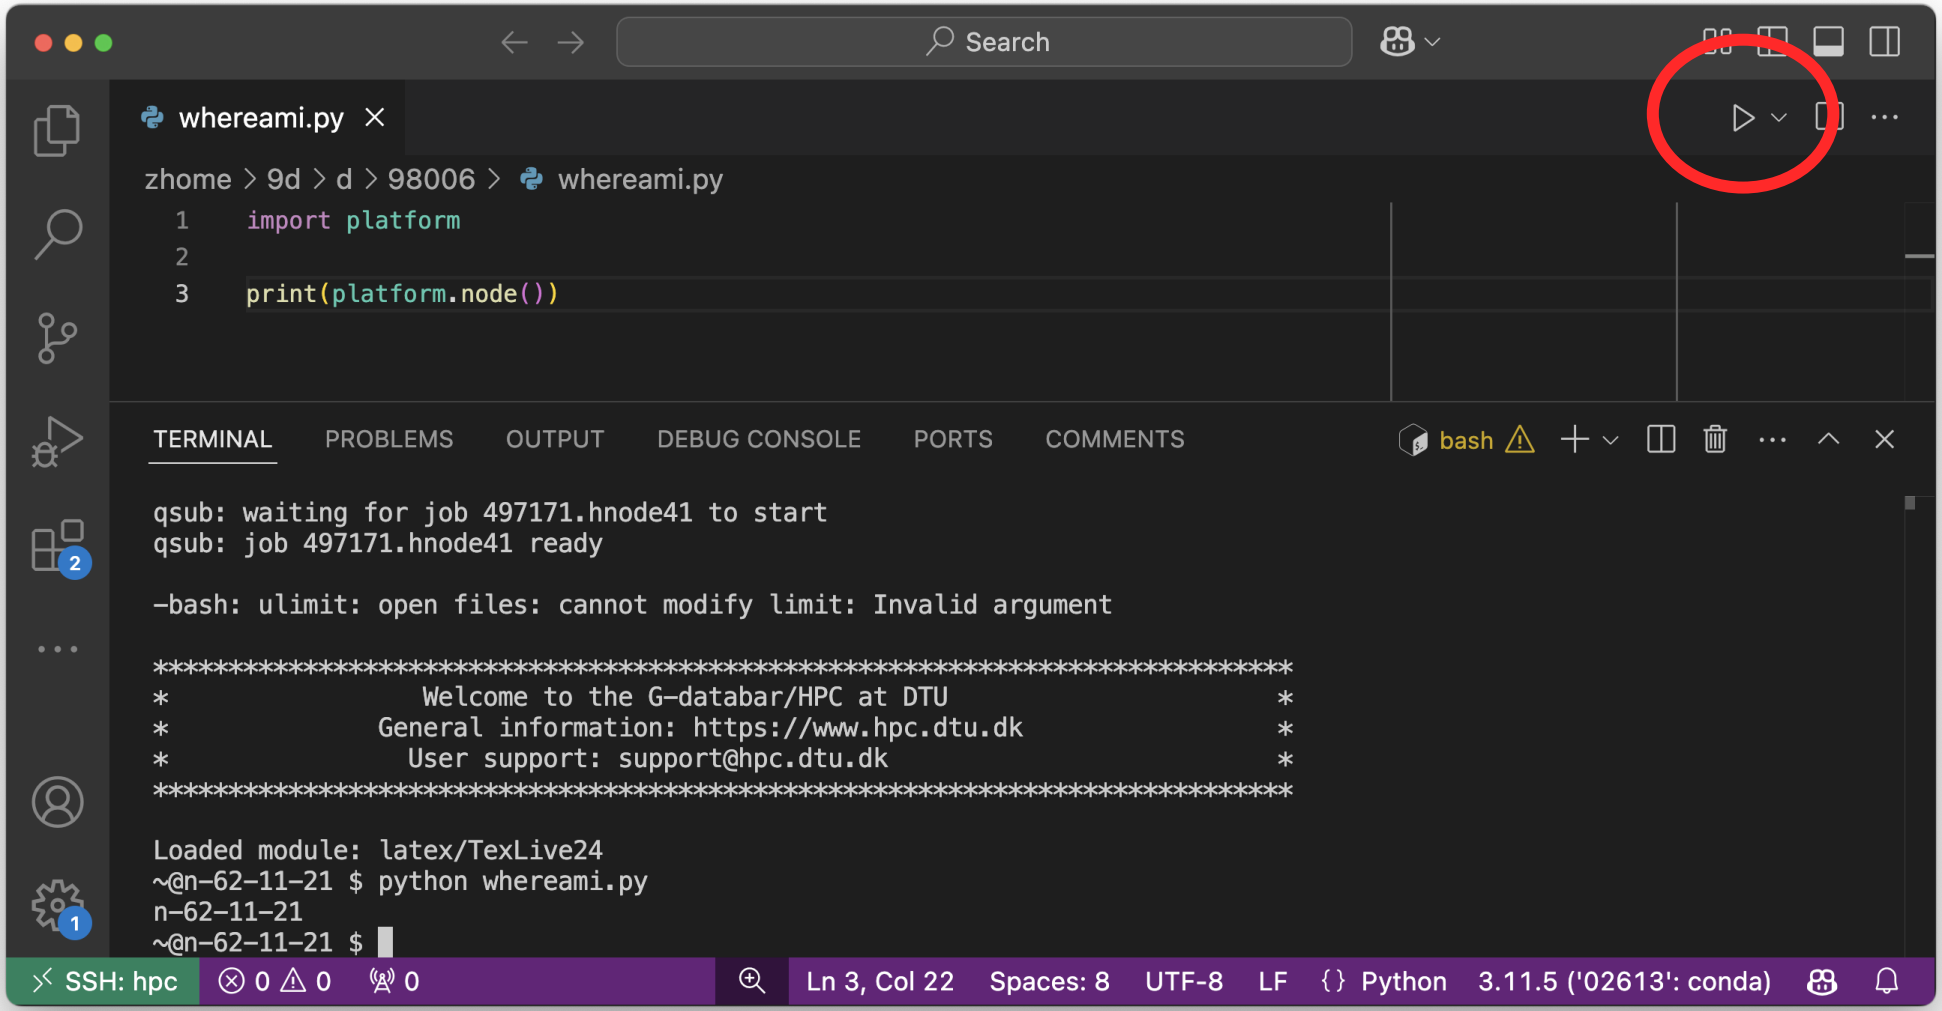

Login nodes and VS Code

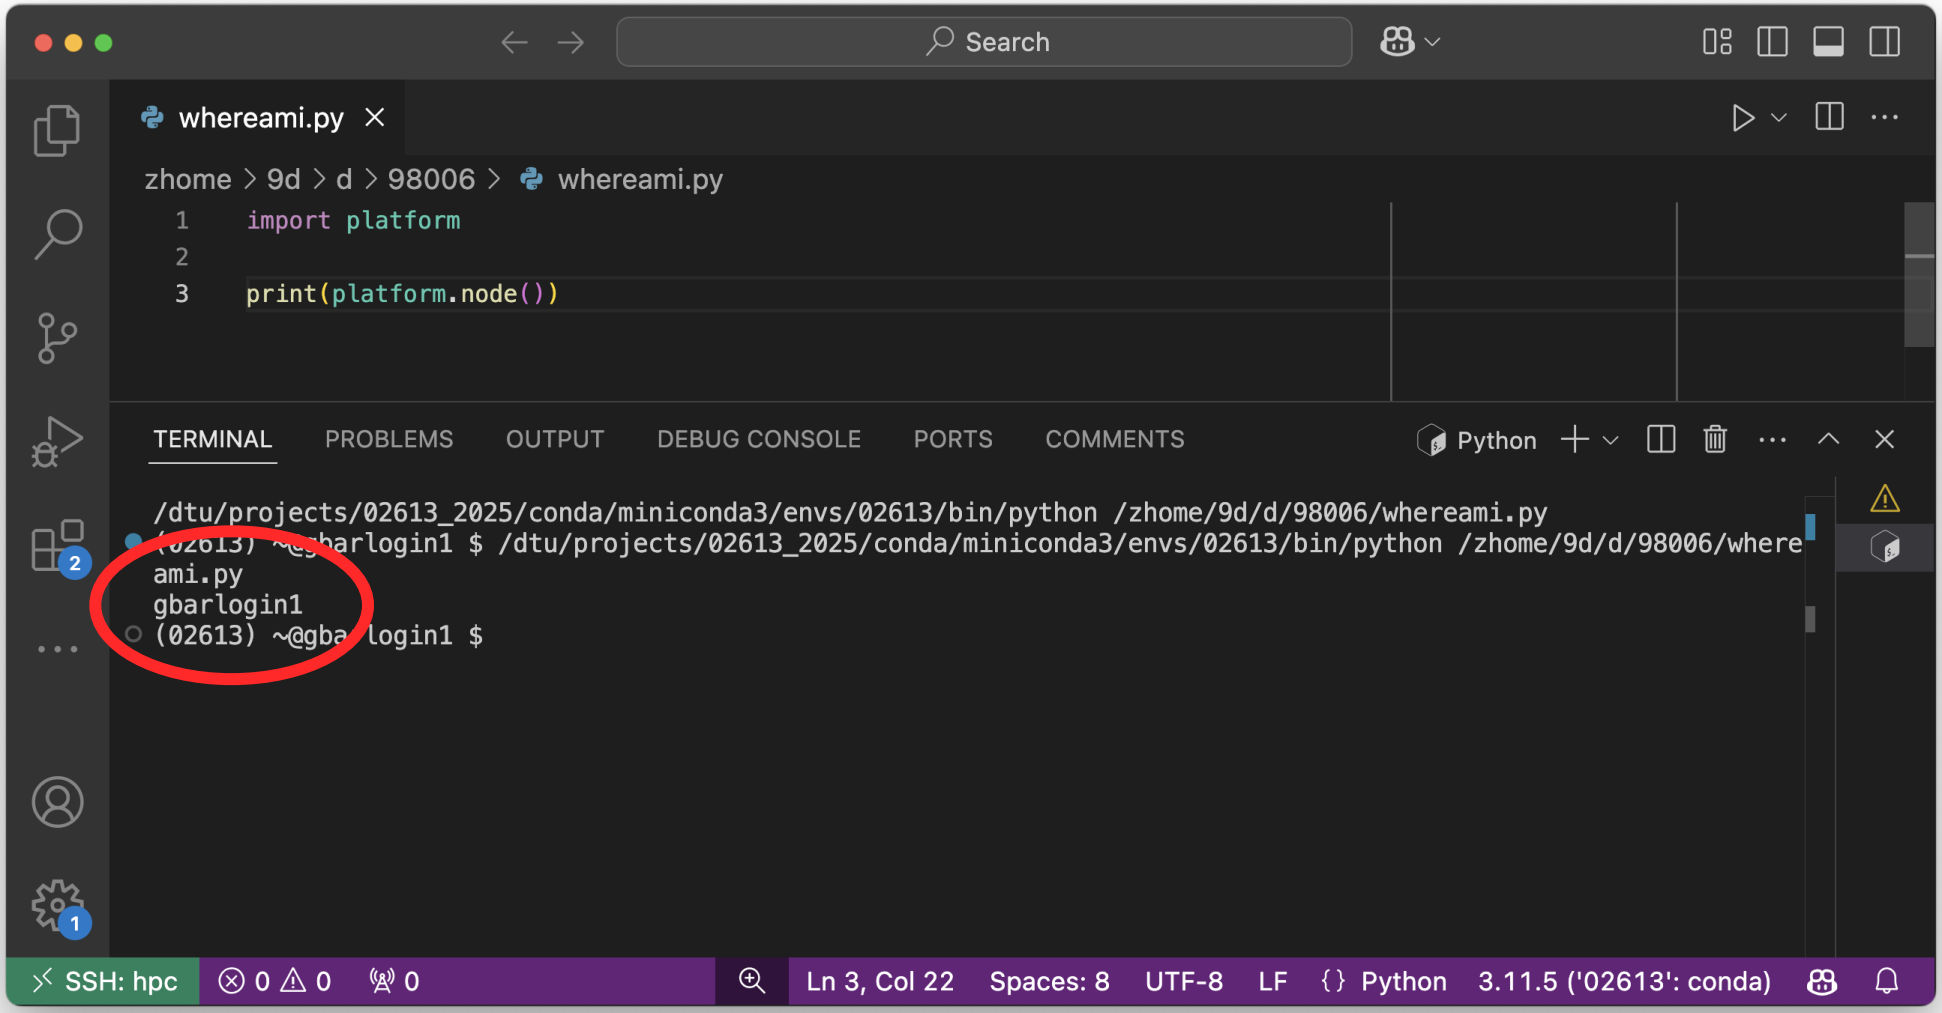

Login nodes and VS Code

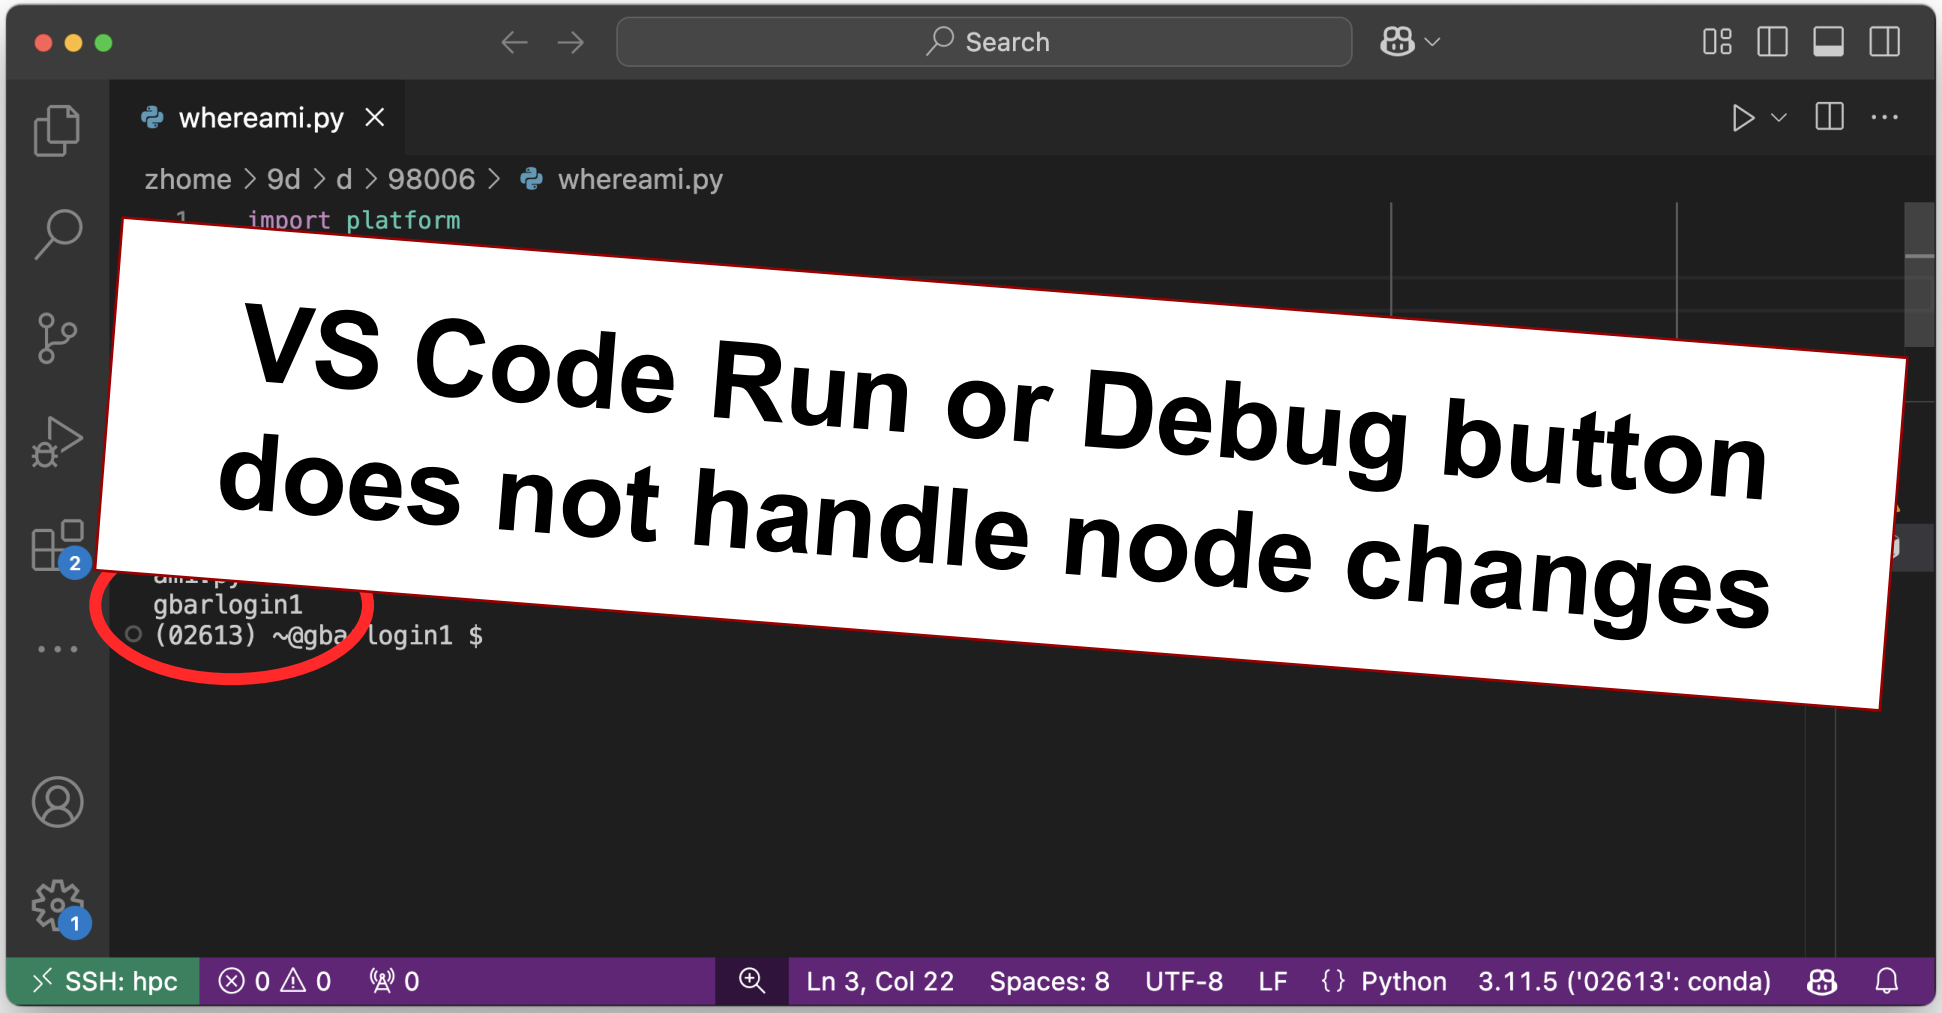

Computer Terminology

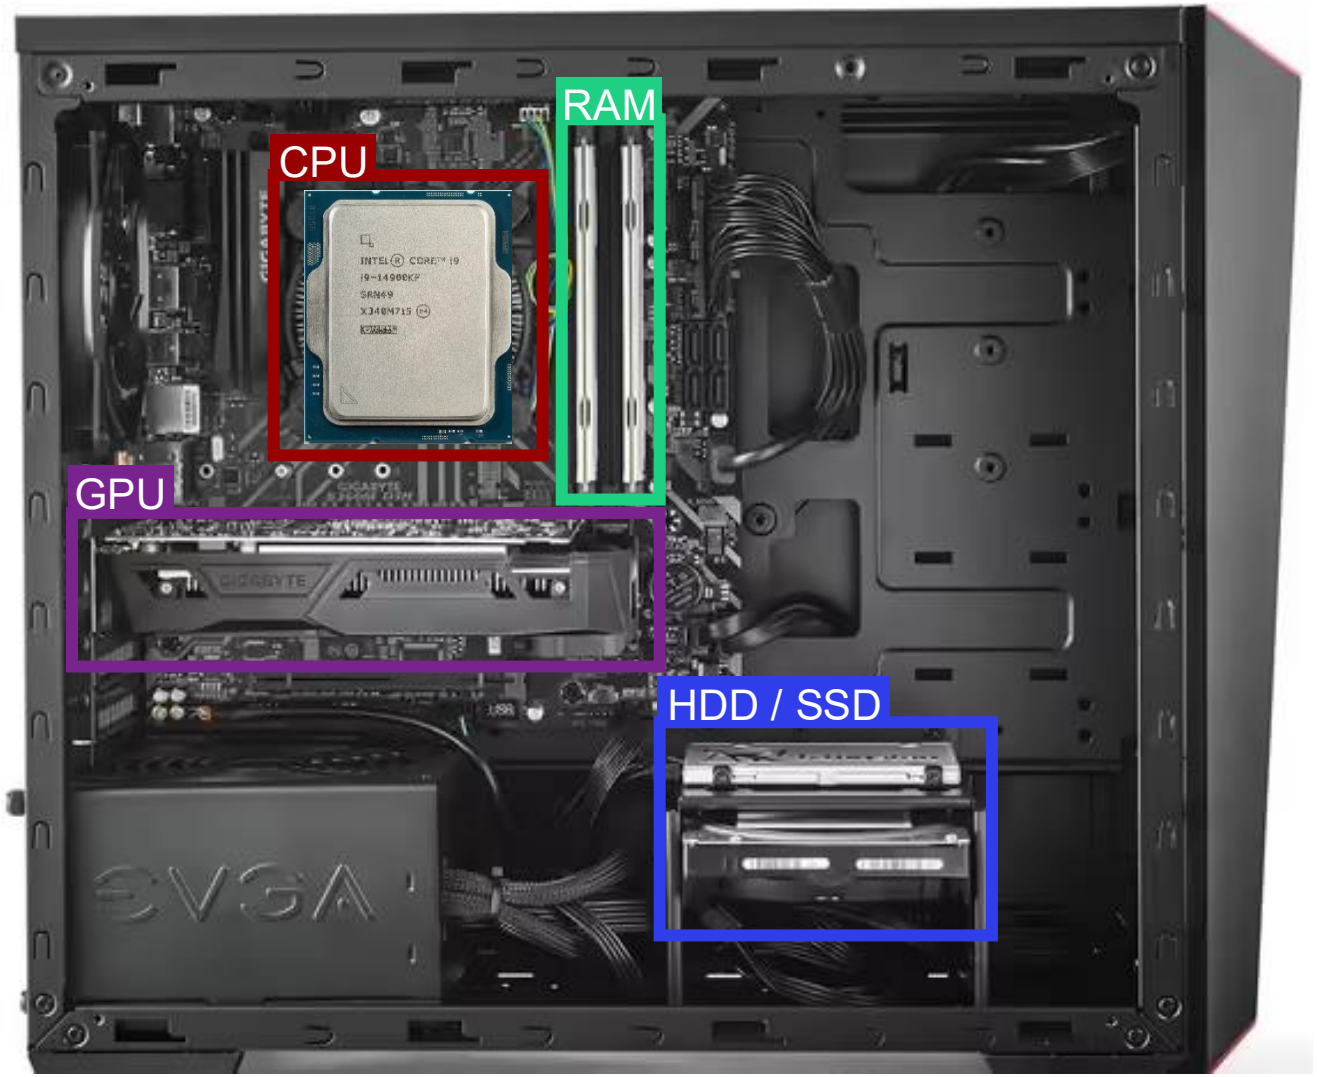

CPU

Performs computations

HDD / SSD

Permanent file storage

RAM

Temporary working memory

Where the CPU “works from”

GPU

Stay tuned for week 9 ;)

Computer Terminology

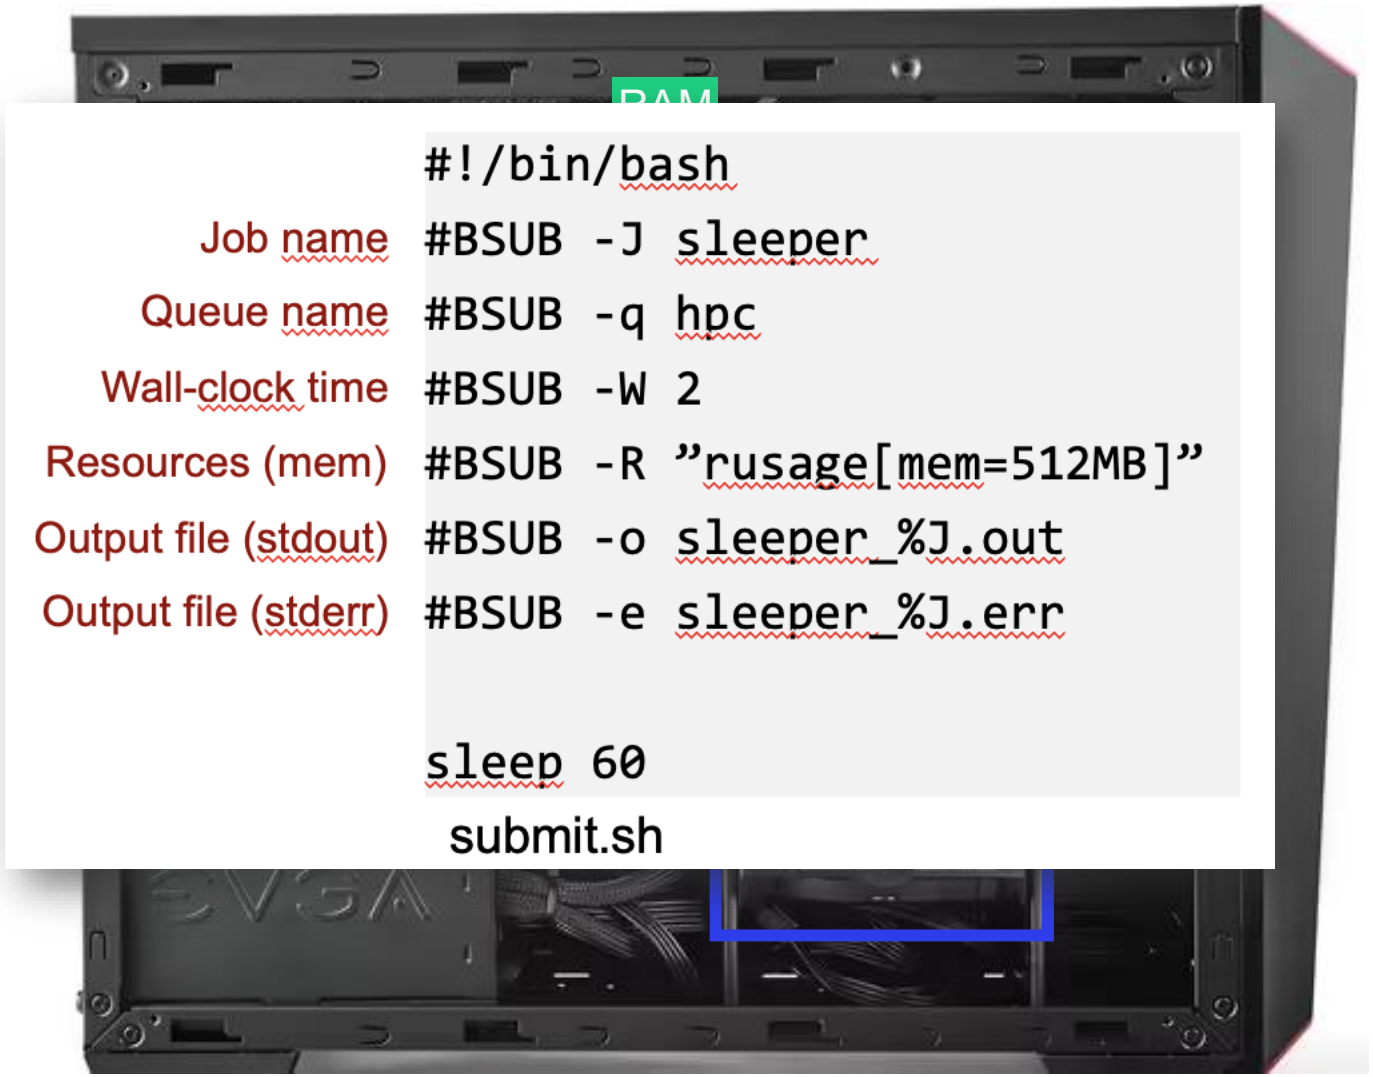

CPU

Performs computations

HDD / SSD

Permanent file storage

RAM

Temporary working memory

Where the CPU "works from"

GPU

Stay tuned for week 9 ;)

Computer Terminology

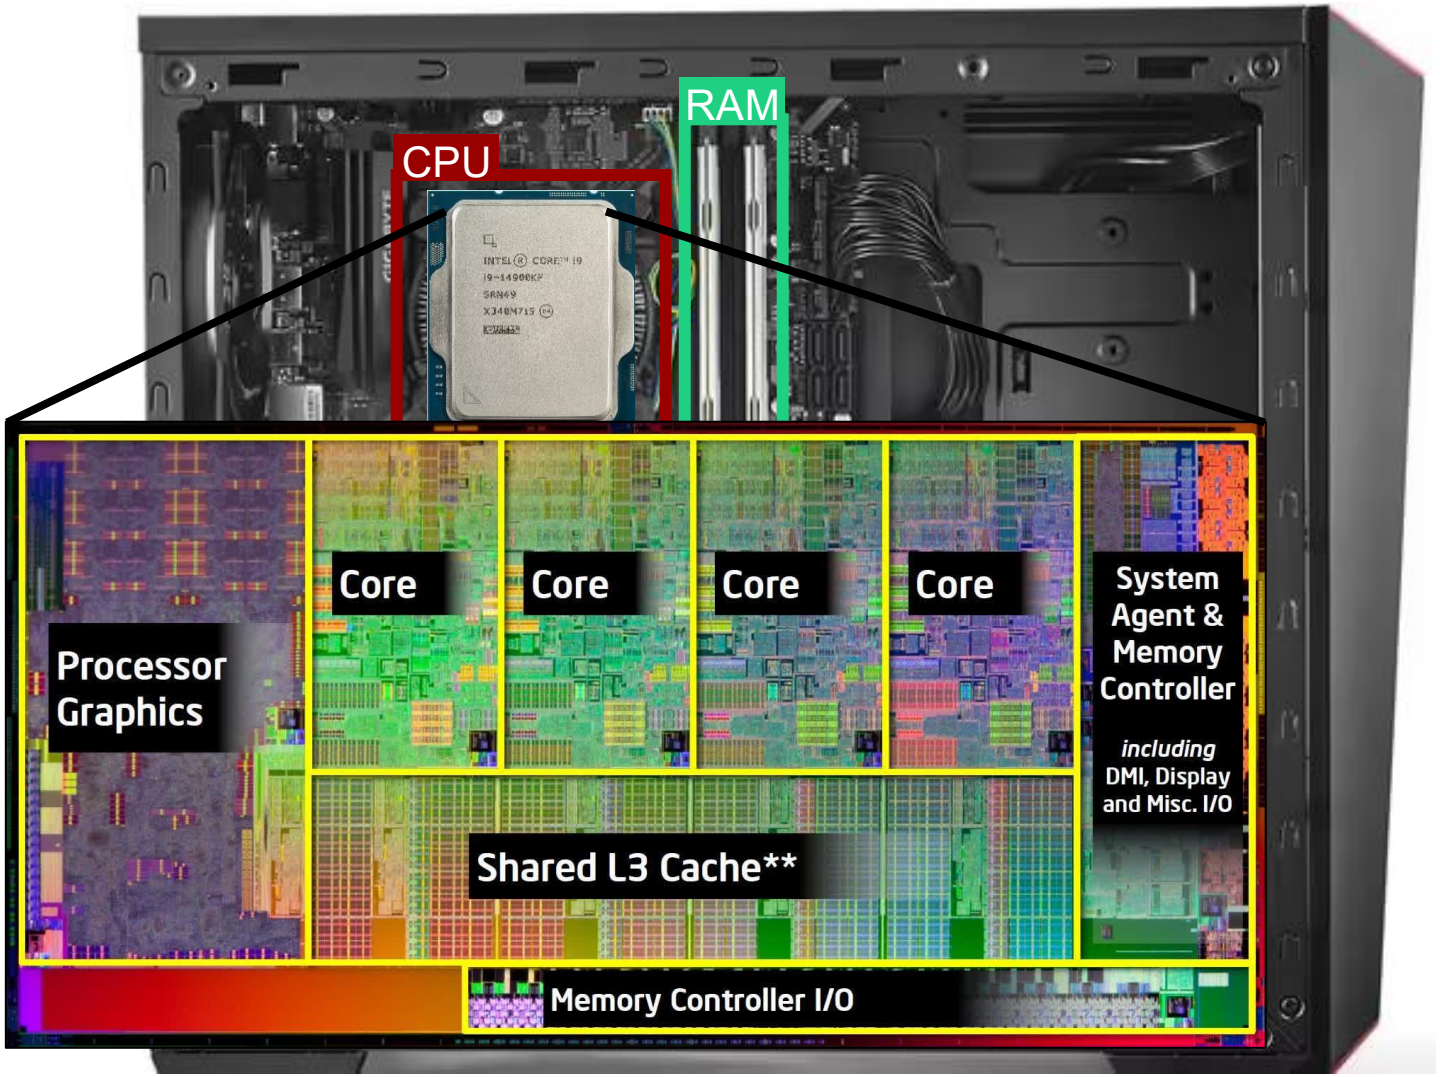

CPU

Performs computations

Contains multiple **cores**

Core can work independently

HDD / SSD

Permanent file storage

RAM

Temporary working memory

Where the CPU “works from”

GPU

Stay tuned for week 9 ;)

Computer Terminology

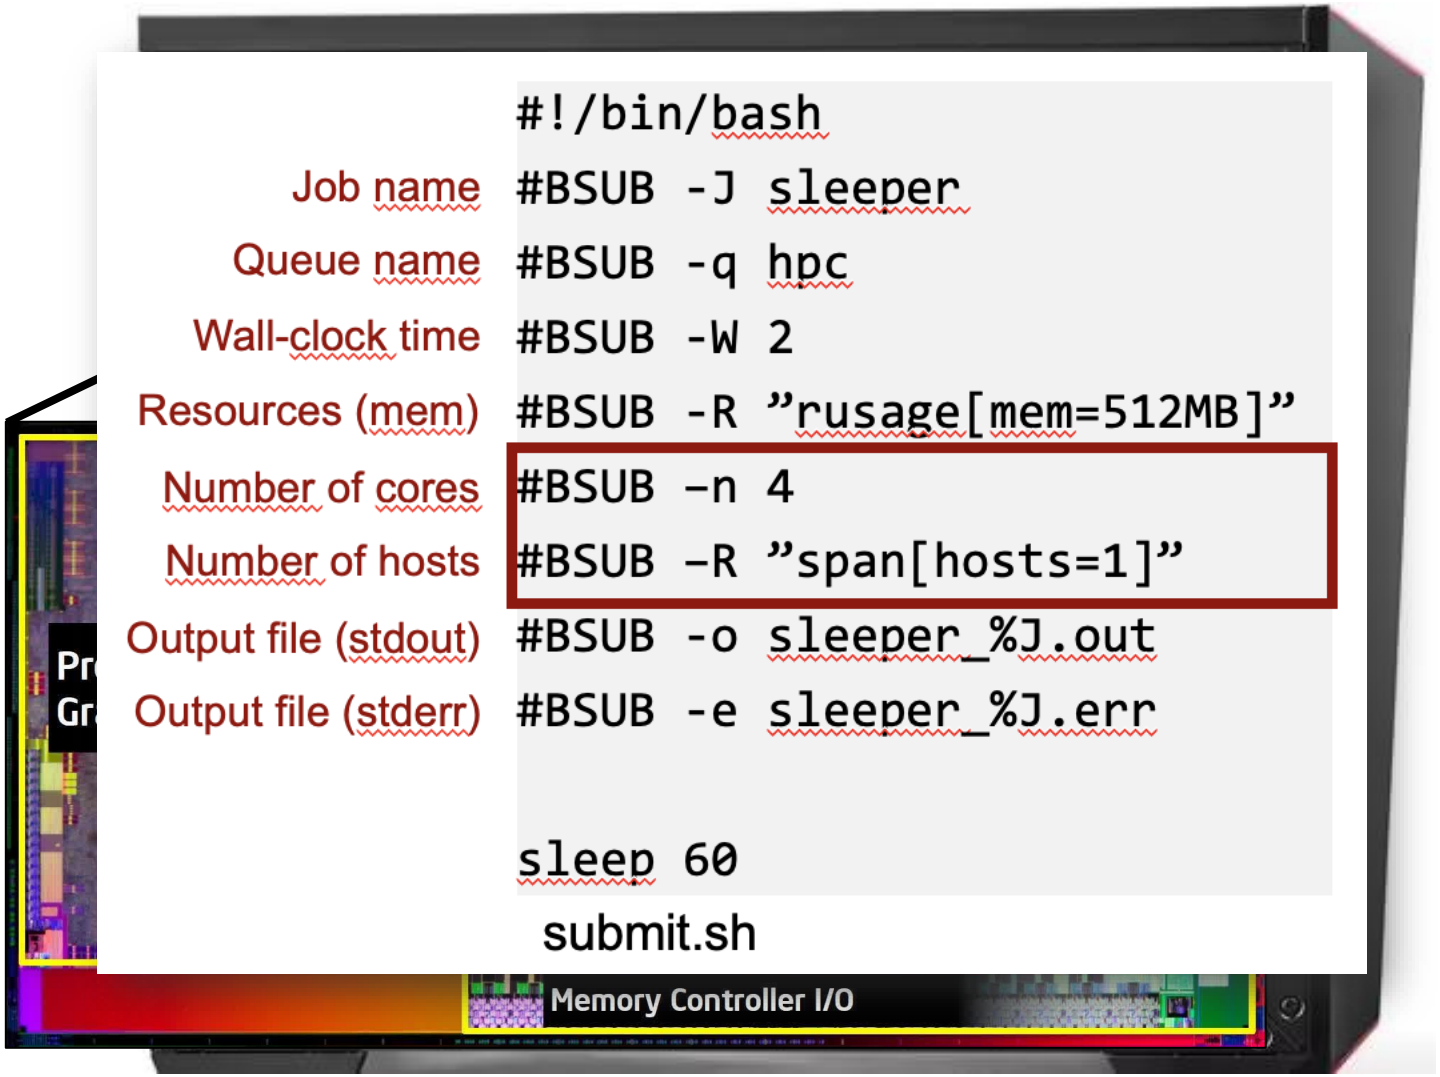

CPU

Performs computations

Contains multiple **cores**

Core can work independently

HDD / SSD

Permanent file storage

RAM

Temporary working memory

Where the CPU "works from"

GPU

Stay tuned for week 9 ;)

Today

1. Why Python?

2. Python programs & command line args

3. Python as batch jobs

Indeed, why Python?

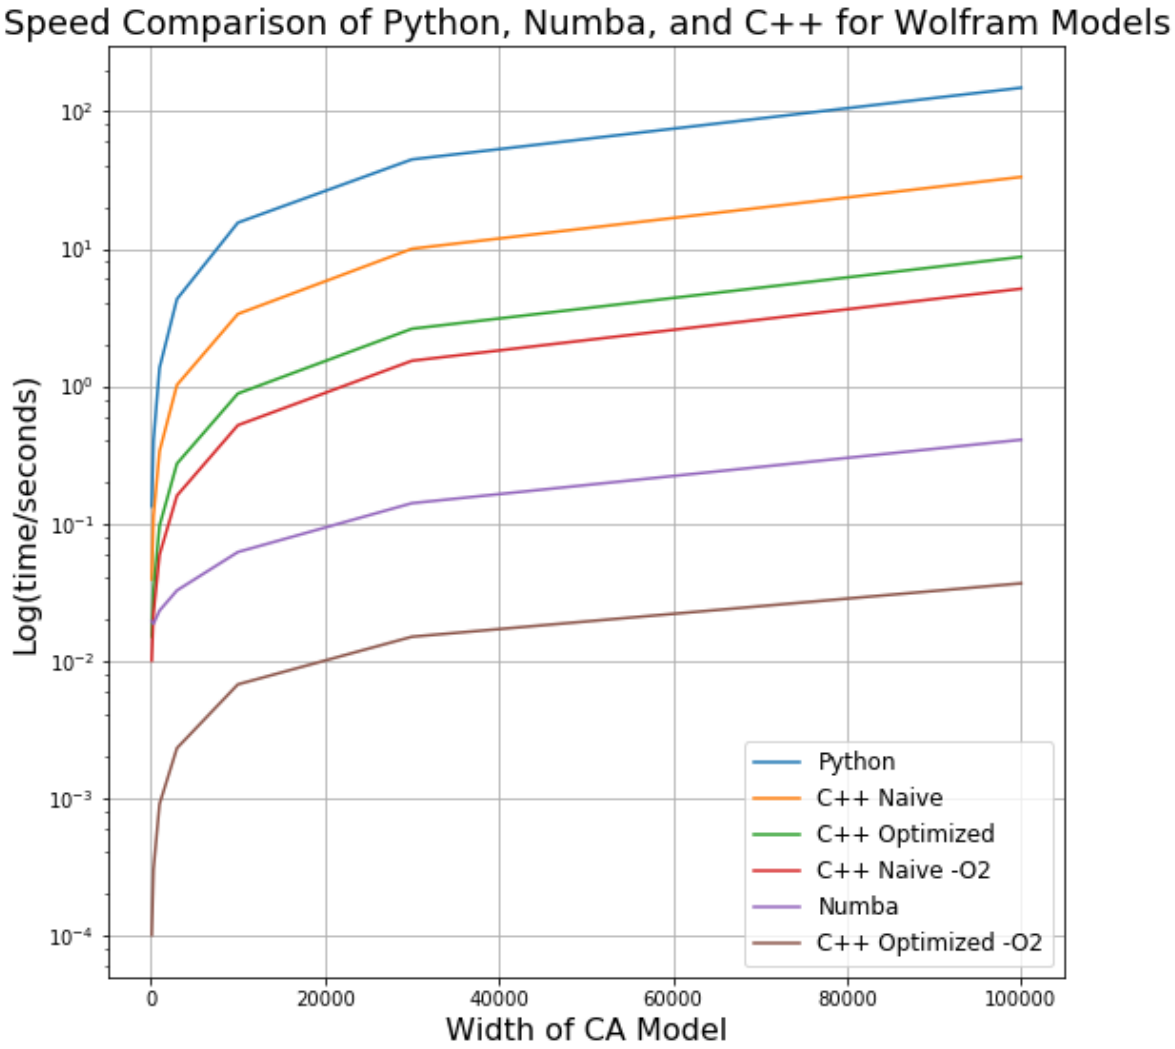

Indeed, why Python?

		fannkuch-redux	n-body	spectral-norm	mandelbrot	pidigits	regex-redux	fasta
C	Time	2.10	2.08	0.40	1.30	0.81	0.80	0.78
	Memory	11,236	11,392	11,392	32,992	11,392	152,148	11,284
Python	Time	285.20	383.12	78.36	155.28	1.25	1.38	49.99
	Memory	14,264	11,096	15,704	14,244	13,564	168,144	11,088

https://benchmarksgame-team.pages.debian.net/benchmarksgame/fastest/python3-gcc.html

Indeed, why Python?

GIL

Global Interpreter Lock

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

What does this print?

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

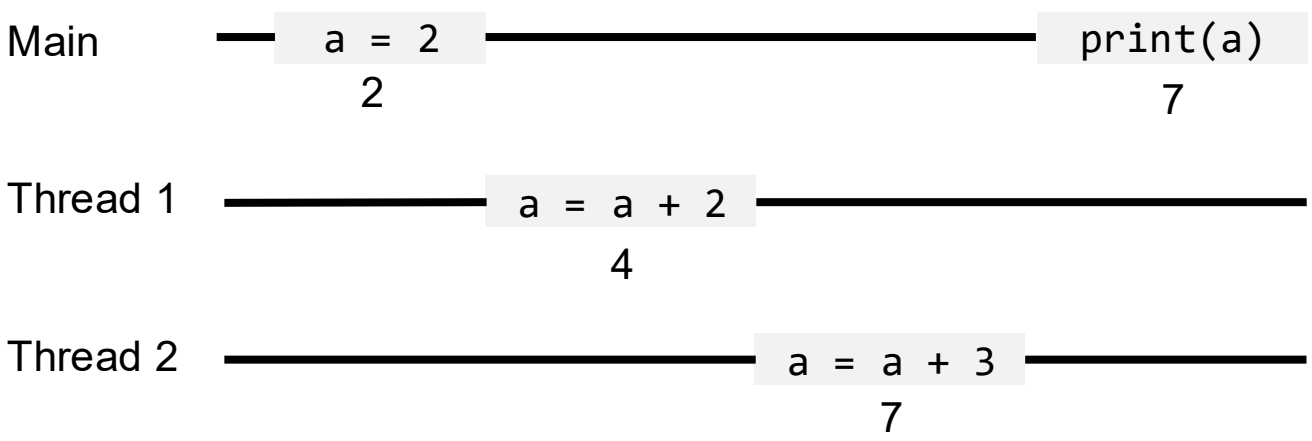

What does this print?

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

What does this print?

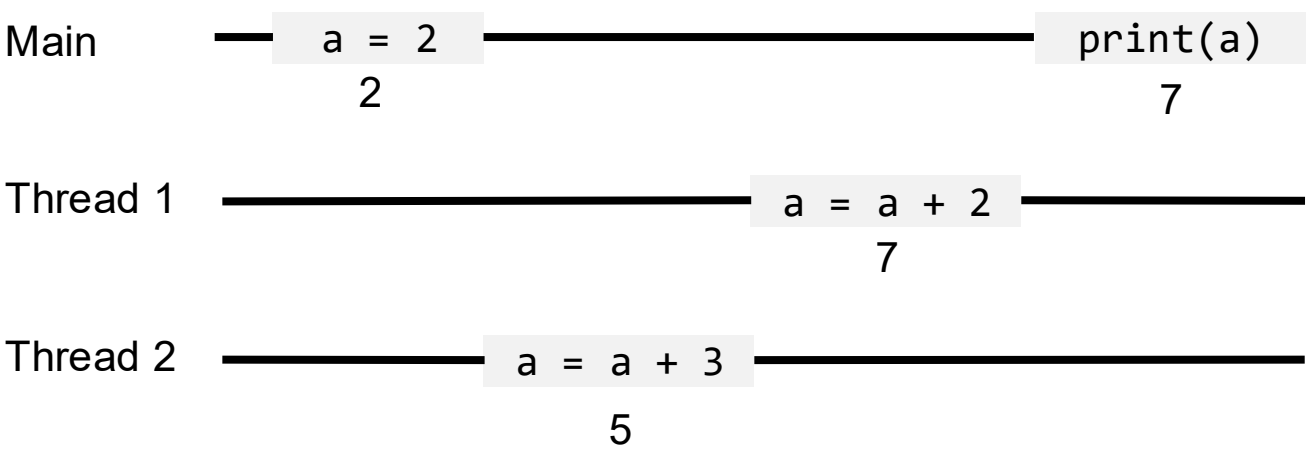

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

What does this print?

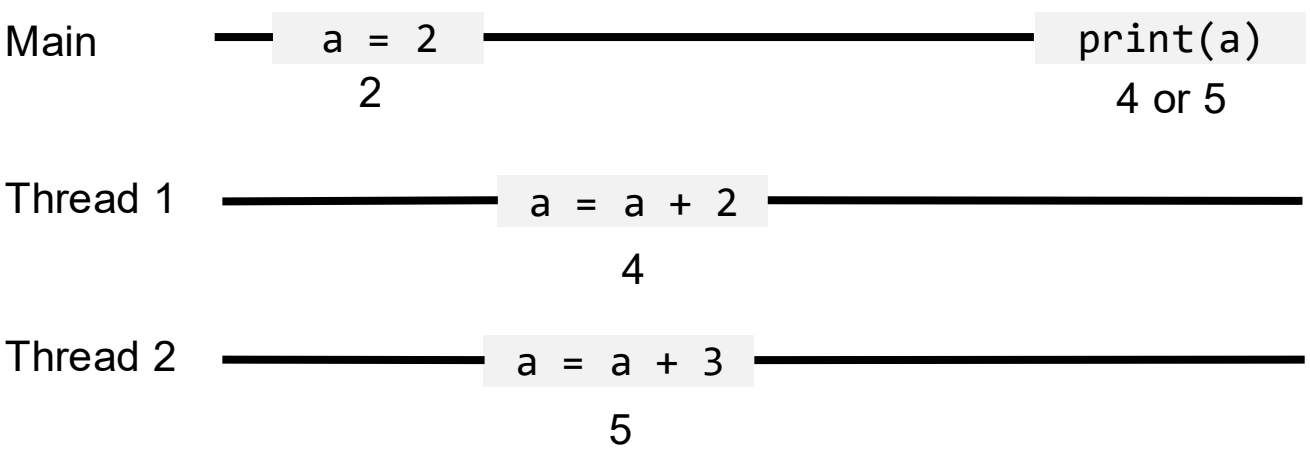

Race condition!

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

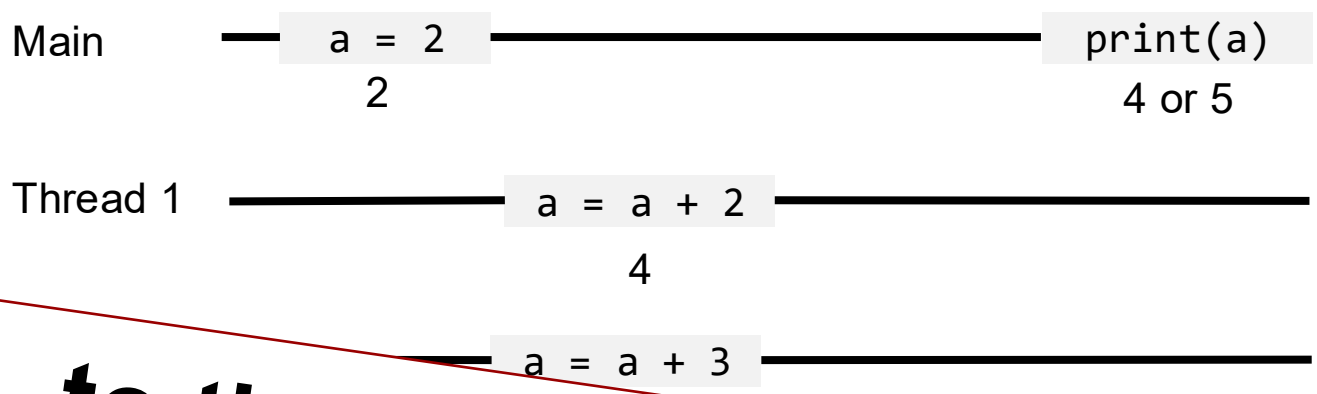

GIL to the rescue!

Race condition!

What does this print?

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

```
<< Main >>
a = 2

<< Thread 1 >>
a = a + 2

<< Thread 2 >>
a = a + 3

<< Main >>
wait_for_threads()
print(a)
```

What does this print?

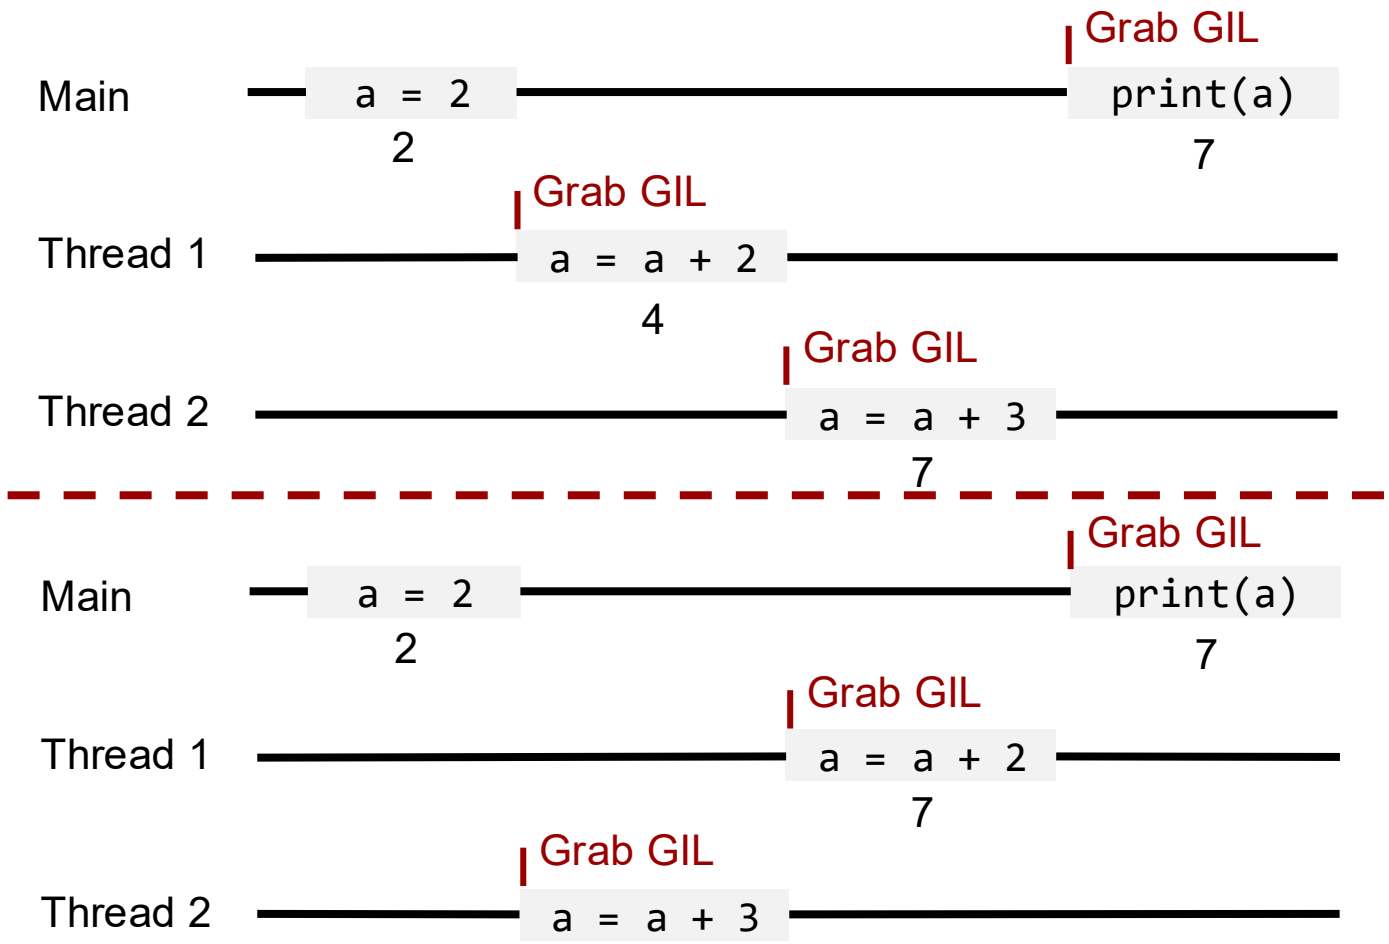

Example from: https://python.land/python-concurrency/the-python-gil

GIL: Global Interpreter Lock

Pro:

- Ensures no race conditions – convenient!
- Some ops release GIL (e.g., I/O, NumPy)

Con:

- Compute heavy multi-threaded = impossible
- Finicky to code around

Instead:

- Use multi-processing
  - more heavy duty
  - harder to share data
- Escape to other languages

GIL: Global Interpreter Lock

Pro:

- Ensures no race conditions – convenient!
- So[illegible] GIL (e.g., I/O, NumPy)

Con:

- Compute heavy multi-threaded = impossible
- Finicky to code around

We'll get back to this in weeks 6 + 7

Instead:

- Use multi-processing
  - more heavy duty
  - harder to share data
- Escape to other languages

Indeed, why Python?

So, Python is

1. Slow
2. Not great for parallel

Python is popular!

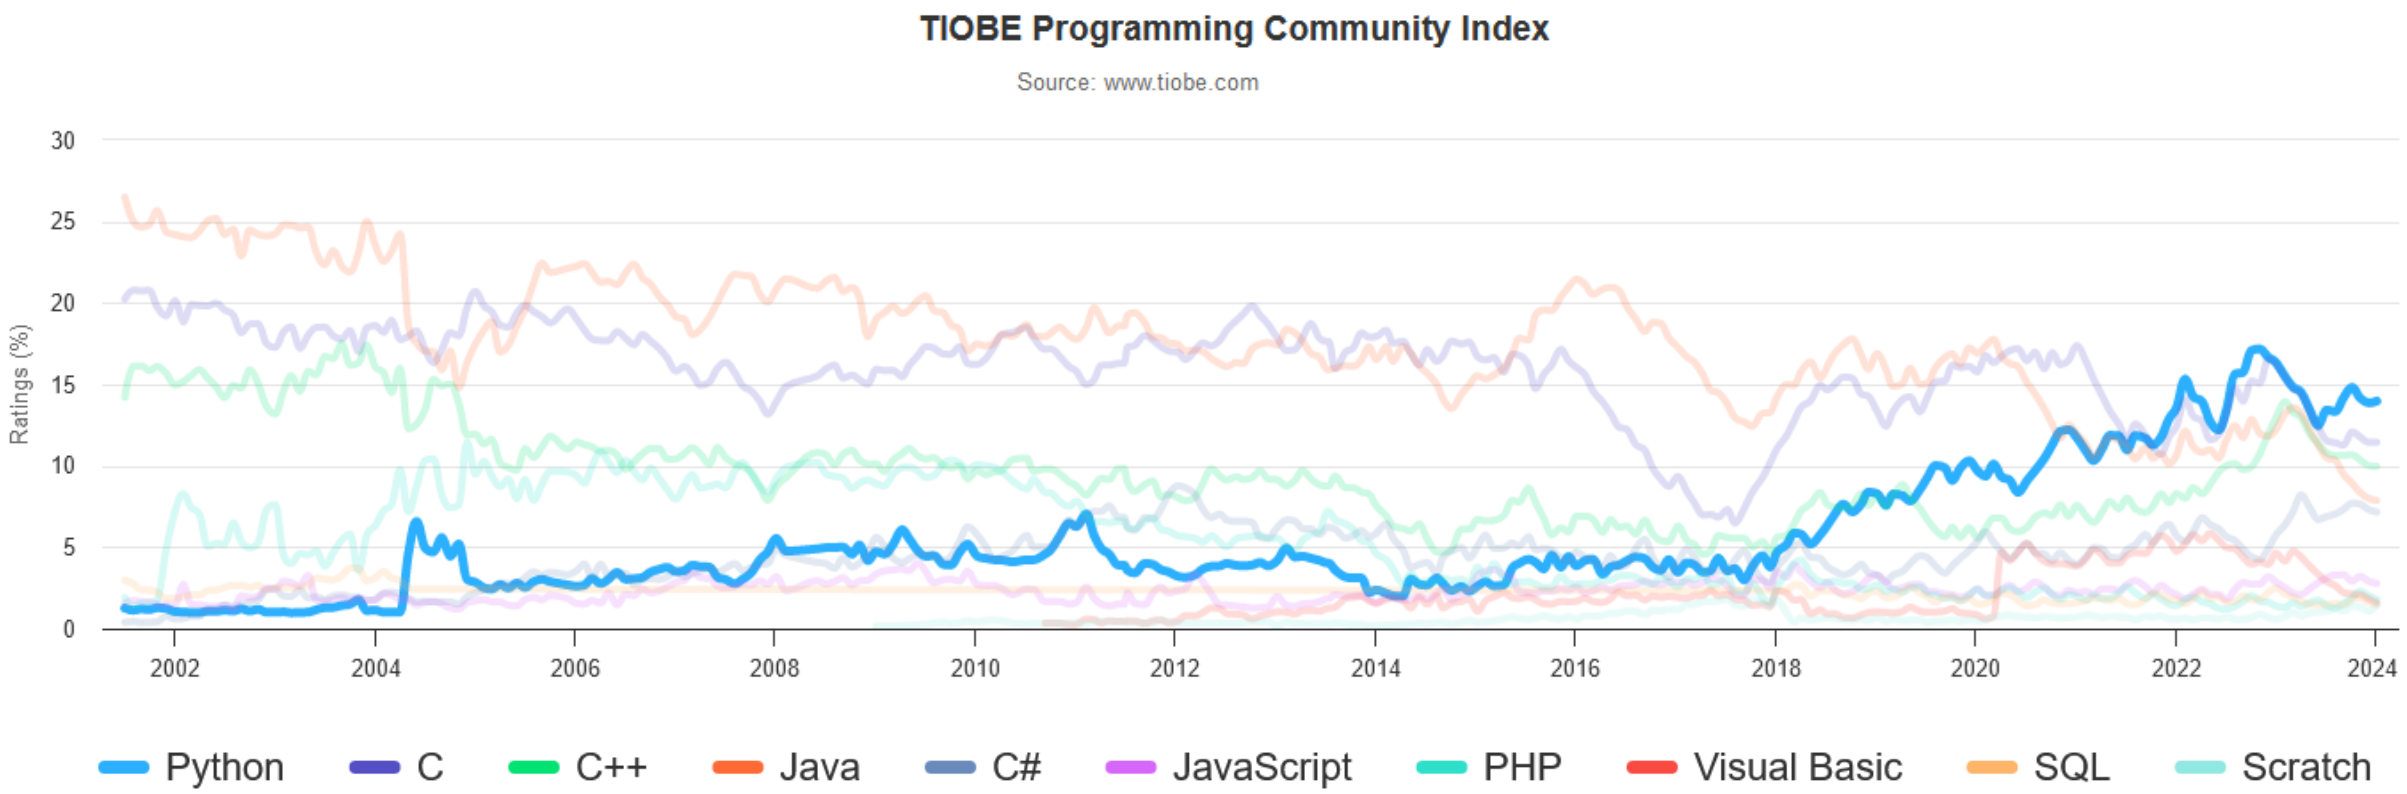

With a great ecosystem

And easy to use

speed vs agility

…and the GIL is going away too

peps.python.org/pep-0703/

PEP 703 – Making the Global Interpreter Lock

- **PEP 703**: CPython 3.13 has experimental support for running with the global interpreter lock disabled.
  See Free-threaded CPython for more details.

Author: Sam Gross <colesbury at gmail.com>
Sponsor: Łukasz Langa <lukasz at python.org>
Discussions-To: Discourse thread
Status: Accepted
Type: Standards Track
Created: 09-Jan-2023
Python-Version: 3.13
Post-History: 09-Jan-2023, 04-May-2023
Resolution: Discourse thread

Join the Vevox session

Go to **vevox.app**

Enter the session ID: **136-979-360**

Or scan the QR code

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

Python is slow because...

it is compiled to machine code

##.##%

it is interpreted

##.##%

snakes have no legs

##.##%

it cannot run in parallel

##.##%

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

Python is slow because...

it is compiled to machine code ##.##%

it is interpreted ##.##%

snakes have no legs ##.##%

it cannot run in parallel ##.##%

RESULTS SLIDE

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

What does GIL stand for?

Global Interpreter Lock

##.##%

Good Interpreted Language

##.##%

Global International Language

##.##%

Grand Interpreter Lie

##.##%

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

What does GIL stand for?

Global Interpreter Lock

##.##%

Good Interpreted Language

##.##%

Global International Language

##.##%

Grand Interpreter Lie

##.##%

RESULTS SLIDE

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

What is the problem with the GIL?

It's hard to pronounce

##.##%

It makes all code slower

##.##%

It must be manually turned on/off

##.##%

It prevents efficient multi-threading

##.##%

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

What is the problem with the GIL?

It's hard to pronounce

##.##%

It makes all code slower

##.##%

It must be manually turned on/off

##.##%

It prevents efficient multi-threading

##.##%

RESULTS SLIDE

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

Why might we still use Python despite being slow?

Snakes are cool!

##.##%

Ease of use increases productivity

##.##%

Nothing else can run on the HPC system

##.##%

Slow code means more breaks for me

##.##%

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

Why might we still use Python despite being slow?

Snakes are cool!

##.##%

Ease of use increases productivity

##.##%

Nothing else can run on the HPC system

##.##%

Slow code means more breaks for me

##.##%

RESULTS SLIDE

Python programs a.k.a. scripts

Python programs

```
print(”Hello World”)
```

hello.py

```
$
```

Python programs

```
print(”Hello World”)
```

hello.py

```
$ python hello.py
Hello World
$
```

Python programs

```
print(”Hello World”)
```

hello.py

```
$ source /dtu/projects/02613_2025/...
$ conda activate 02613
(02613) $ python hello.py
Hello World
(02613) $
```

Python programs

```
print(”Hello World”)
```

hello.py

```
$ python hello.py
Hello World
$
```

Python programs

```
name = ”02613”
print(f”Hello {name}”)
```

hello.py

```
$ python hello.py
Hello 02613
$
```

Python programs

```
name = ??
print(f”Hello {name}”)
```

hello.py

```
$ python hello.py 02613
Hello 02613
$
```

Python programs

```
def compute_from_data(file_path):
    data = load_data(file_path)
    # Compute...
    ...
    return result



result = compute_from_data(”path/to/data.txt”)
print(result)
```

program.py

```
$
```

Python programs

```
def compute_from_data(file_path):
    data = load_data(file_path)
    # Compute...
    ...
    return result



result = compute_from_data(”path/to/data.txt”)
print(result)
```

program.py

```
$ python program.py
42
$
```

Python programs

```
def compute_from_data(file_path):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = ???
result = compute_from_data(file_path)
print(result)
```

program.py

```
$ python program.py path/to/data.txt
42
$
```

Python programs

```
import sys

print(sys.argv)
```

List of command line arguments:
0: Script file name
1: First argument
2: Second argument
…

```
$
```

program.py

Python programs

```
import sys

print(sys.argv)
```

List of command line arguments:
0: Script file name
1: First argument
2: Second argument
…

program.py

```
$ python program.py path/to/data.txt
['program.py', 'path/to/data.txt']
$
```

Python programs

```
import sys

print(sys.argv)
```

List of command line arguments:
0: Script file name
1: First argument
2: Second argument
…

program.py

```
$ python program.py path/to/data.txt
['program.py', 'path/to/data.txt']
$ python program.py 1st 2nd 3rd
['program.py', '1st', '2nd', '3rd']
$
```

Python programs

```
import sys

print(sys.argv)
```

List of command line arguments:

0: Script file name

1: First argument

2: Second argument

…

program.py

```
$ python program.py path/to/data.txt
['program.py', 'path/to/data.txt']
$ python program.py 1st 2nd 3rd
['program.py', '1st', '2nd', '3rd']
$ python program.py
```

Python programs

```
import sys

print(sys.argv)
```

List of command line arguments:

0: Script file name

1: First argument

2: Second argument

…

program.py

```
$ python program.py path/to/data.txt
['program.py', 'path/to/data.txt']
$ python program.py 1st 2nd 3rd
['program.py', '1st', '2nd', '3rd']
$ python program.py
['program.py']
$
```

Python programs

```
def compute_from_data(file_path):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = ???
result = compute_from_data(file_path)
print(result)
```

program.py

```
$ python program.py path/to/data.txt
42
$
```

Python programs

```python
import sys

def compute_from_data(file_path):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = sys.argv[1]
result = compute_from_data(file_path)
print(result)
```

program.py

```
$ python program.py path/to/data.txt
42
$
```

Python programs

```
import sys

def compute_from_data(file_path, num_iters):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = sys.argv[1]
num_iters = sys.argv[2]
result = compute_from_data(file_path, num_iters)
print(result)
```

program.py

```
$ python program.py path/to/data.txt 100
```

Python programs

```
import sys

def compute_from_data(file_path, num_iters):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = sys.argv[1]
num_iters = sys.argv[2]
result = compute_from_data(file_path, num_iters)
print(result)
```

program.py

```
$ python program.py path/to/data.txt 100
Traceback (most recent call last):
…
TypeError: 'str' object cannot be
interpreted as an integer
$
```

Python programs

```python
import sys

def compute_from_data(file_path, num_iters, a):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = sys.argv[1]
num_iters = int(sys.argv[2])  # Cast to int
result = compute_from_data(file_path, num_iters, a)
print(result)
```

program.py

```
$ python program.py path/to/data.txt 100
42
$
```

Python programs

```
import sys

def compute_from_data(file_path, num_iters, a):
    data = load_data(file_path)
    # Compute...
    ...
    return result


file_path = sys.argv[1]
num_iters = int(sys.argv[2])  # Cast to int
a = float(sys.argv[3])        # Cast to float
result = compute_from_data(file_path, num_iters, a)
print(result)
```

program.py

```
$ python program.py path/to/data.txt 100 0.2
42
$
```

Join the Vevox session

Go to **vevox.app**

Enter the session ID: **136-979-360**

Or scan the QR code

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

How do you access command line arguments?

A ##.##%

B ##.##%

C ##.##%

D ##.##%

A

```
import sys
sys.args
```

B

```
import cli
cli.args
```

C

```
import sys
sys.argv
```

D

```
argv = input()
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

How do you access command line arguments?

A ##.##%

B ##.##%

C ##.##%

D ##.##%

A

```
import sys
sys.args
```

B

```
import cli
cli.args
```

C

```
import sys
sys.argv
```

D

```
argv = input()
```

Join at: **vevox.app**

ID: **XXX-XXX-XXX**

Question slide

What will this print?

A

##.##%

B

##.##%

C

##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py hello
```

A

```
['python', 'prog.py', 'hello']
```

B

```
['prog.py', 'hello']
```

C

```
['hello']
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

What will this print?

A ##.##%

B ##.##%

C ##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py hello
```

A

```
['python', 'prog.py', 'hello']
```

B

```
['prog.py', 'hello']
```

C

```
['hello']
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

What will this print?

A

##.##%

B

##.##%

C

##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py data/my data.d
```

A

```
['prog.py', 'data/my data.d']
```

B

```
['prog.py', 'data/my', 'data.d']
```

C

```
['prog.py', 'data', 'my data.d']
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

What will this print?

A

##.##%

B

##.##%

C

##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py data/my data.d
```

A

```
['prog.py', 'data/my data.d']
```

B

```
['prog.py', 'data/my', 'data.d']
```

C

```
['prog.py', 'data', 'my data.d']
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Question slide

What will this print?

A ##.##%

B ##.##%

C ##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py x 1 2.0
```

A

```
['prog.py', 'x', 1, 2.0]
```

B

```
['prog.py', 'x', '12.0']
```

C

```
['prog.py', 'x', '1', '2.0']
```

Join at: **vevox.app** ID: **XXX-XXX-XXX**

Results slide

What will this print?

A ##.##%

B ##.##%

C ##.##%

prog.py:

```
import sys
print(sys.args)
```

```
$ python prog.py x 1 2.0
```

A

```
['prog.py', 'x', 1, 2.0]
```

B

```
['prog.py', 'x', '12.0']
```

C

```
['prog.py', 'x', '1', '2.0']
```

Python in batch jobs

Python in batch jobs

Resources

Command

```
#!/bin/bash
#BSUB -J sleeper
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o sleeper_%J.out
#BSUB -e sleeper_%J.err

sleep 60
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs

Resources

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs

Resources

Initialize

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs

Resources

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
```

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Your job script is read at submission, but…
Your Python script is read when the job starts!

Run

```
# Run Python script
python program.py
```

submit.sh

Python in batch jobs

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py path/to/data.txt 0.2
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

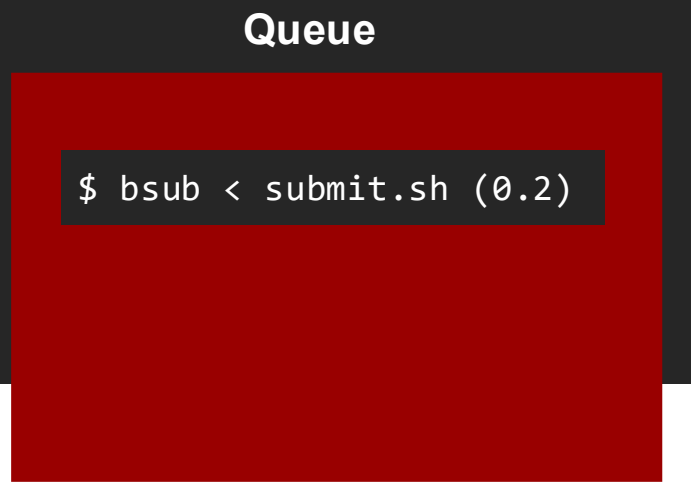

Python in batch jobs

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py path/to/data.txt 0.4
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bsub < submit.sh
Job <702574> is submitted to queue <hpc>.
```

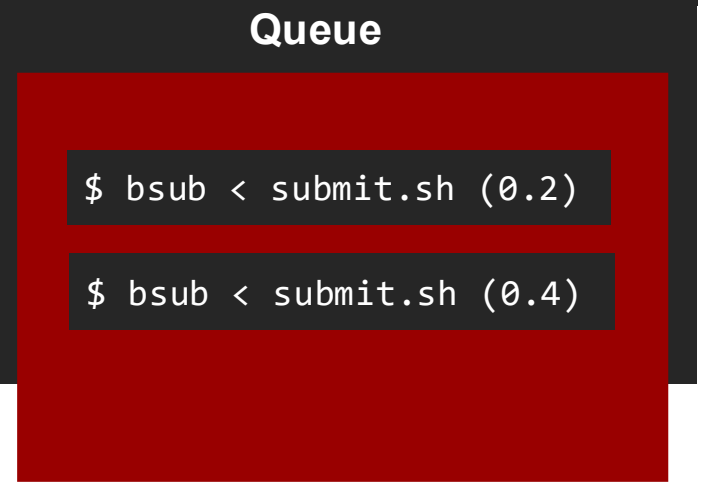

Python in batch jobs

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py path/to/data.txt 0.6
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bsub < submit.sh
Job <702574> is submitted to queue <hpc>.
$ bsub < submit.sh
Job <702575> is submitted to queue <hpc>.
```

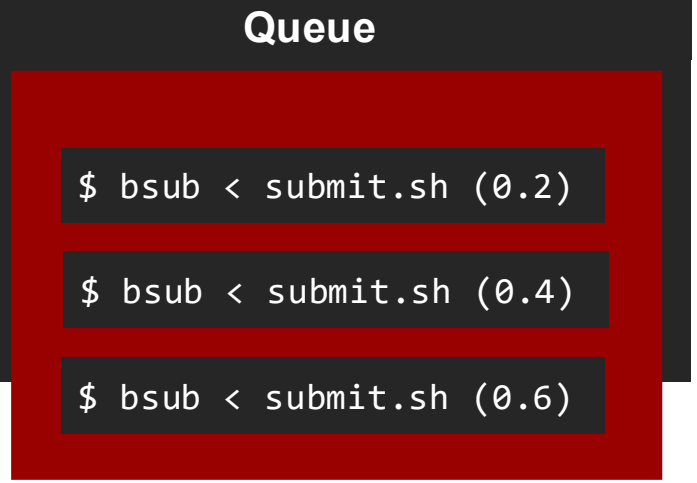

Python in batch jobs: some tips

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bstat
JOBID  ... STAT START_TIME    ELAPSED
702573 ... RUN  Dec 14 13:48  0:04:32
$
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

No ”Hello World”!

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bstat
JOBID  ... STAT START_TIME    ELAPSED
702573 ... RUN  Dec 14 13:48  0:04:32
$ bpeek
<< output from stdout >>

<< output from stderr >>
$
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

No ”Hello World”!

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bstat
JOBID  ... STAT START_TIME    ELAPSED
702573 ... RUN  Dec 14 13:48  0:04:32
$ bpeek
<< output from stdout >>

<< output from stderr >>
$ cat python_702573.out
$
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ bstat
No unfinished job found
$ cat python_702573.out
Hello World

------------------------------------------
Sender: LSF System <lsfadmin@hpc.dtu.dk>
…
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

submit.sh

Unbuffered output

```
$ bsub < submit.sh
Job <702574> is submitted to queue <hpc>.
$ bstat
JOBID  ... STAT START_TIME    ELAPSED
702574 ... RUN  Dec 14 13:48  0:04:32
$ bpeek
<< output from stdout >>
Hello World

<< output from stderr >>
$
```

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: unbuffered output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BS

# I
sou
con

# Run Python script
python -u program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702574> is submitted to queue <hpc>.
$ bstat
JOBID   ... STAT START_TIME    ELAPSED
702574 ... RUN  Dec 14 13:48  0:04:32
```

Has performance implications!
Not good if you print a lot

Unbuffered output

```
import time
print(”Hello World”)
time.sleep(600)  # Sleep 10 minutes
```

program.py

Python in batch jobs: organizing your output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs: organizing your output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o python_%J.out
#BSUB -e python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ ls
program.py
python_702531.err
python_702531.out
python_702572.err
python_702571.err
python_702571.out
python_702572.err
python_702572.out
python_702573.err
python_702573.out
…
```

Python in batch jobs: organizing your output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o batch_output/python_%J.out
#BSUB -e batch_output/python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

Directory for output files

submit.sh

```
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs: organizing your output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[mem=512MB]”
#BSUB -o batch_output/python_%J.out
#BSUB -e batch_output/python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

Directory for output files

submit.sh

Make directory before bsub!

```
$ mkdir batch_output
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$
```

Python in batch jobs: organizing your output

Resources

Initialize

Run

```
#!/bin/bash
#BSUB -J python
#BSUB -q hpc
#BSUB -W 15
#BSUB -R ”rusage[me=512MB]”
#BSUB -o batch_output/python_%J.out
#BSUB -e batch_output/python_%J.err

# Initialize Python environment
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

# Run Python script
python -u program.py
```

Directory for output files

submit.sh

Make directory before bsub!

```
$ mkdir batch_output
$ bsub < submit.sh
Job <702573> is submitted to queue <hpc>.
$ ls
batch_output
program.py
submit.sh
```

Today’s exercise

Get comfortable with Python on HPC

- Run Python interactively for testing
- Touch on relevant Python concepts

  command line args, functions, lambdas, classes, list comprehensions, basic NumPy, basic I/O
- Run Python in batch jobs
- **Remember:** Autolab exercises today do not count towards mandatory 20. But must still complete 1 (one from every week).

Useful commands

Log on to HPC
ssh <username>@login.hpc.dtu.dk
ssh <username>@login2.hpc.dtu.dk
ssh <username>@login.gbar.dtu.dk
ssh <username>@login2.gbar.dtu.dk

Init. Anaconda environment (also for job scripts)
source /dtu/projects/02613_2025/conda/conda_init.sh
conda activate 02613

Python with unbuffered output
python -u program.py

Python command line arguments
import sys
sys.argv # List of arguments

Change to work node
linuxsh

Submit job script
bsub < submit.sh

Job status
bstat
bjobs
bjobs –p # pending reason

Check job output
bpeek
bpeek <JOBID>

Kill job
bkill <JOBID>