Join the Vevox session

Go to **vevox.app**

Enter the session ID: **176-017-158**

Or scan the QR code

What was the strategy for the final GPU reduction?

The kernel performs one layer of reduction

0%

The kernel reduces over one thread block

0%

The kernel performs the whole reduction

0%

What was the strategy for the final GPU reduction?

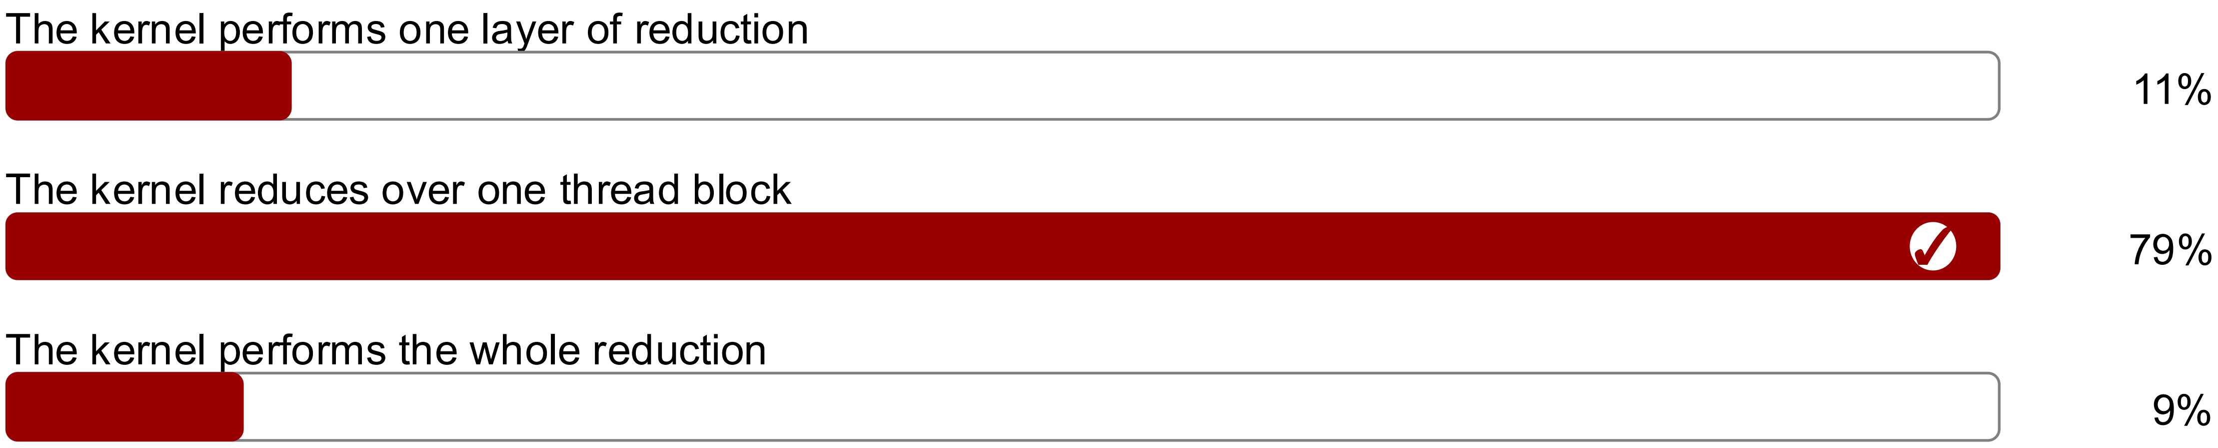

RESULTS SLIDE

What was the purpose of shared memory?

Makes it faster to transfer memory to/from the GPU — 0%

Very fast memory all threads can access — 0%

Very fast memory threads in the same block can access — 0%

What was the purpose of shared memory?

Makes it faster to transfer memory to/from the GPU 12%

Very fast memory all threads can access 19%

Very fast memory threads in the same block can access ✓ 69%

RESULTS SLIDE

When should we use shared memory?

Always

0%

When threads repeatedly read the same memory

0%

When we can't rely on L1 cache

0%

When should we use shared memory?

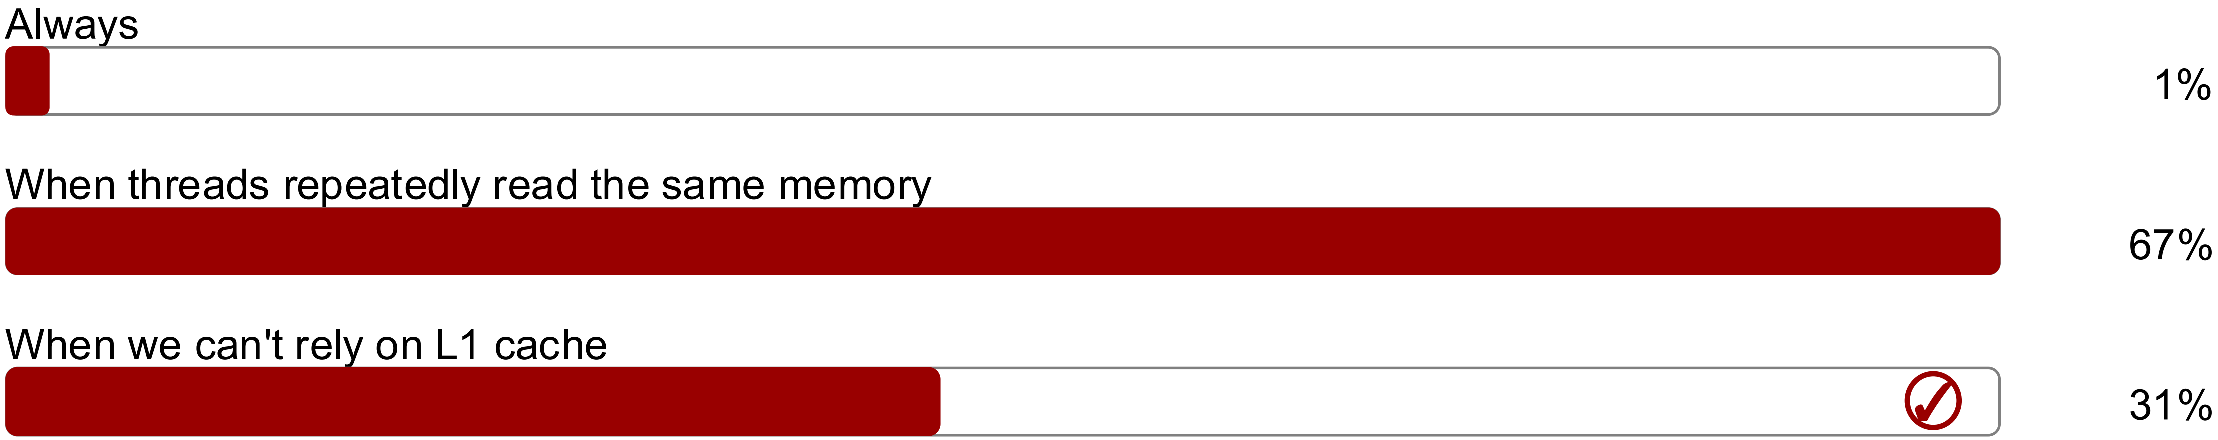

RESULTS SLIDE

Back to the slides